

# Prevention of Microarchitectural Covert Channels on an Open-Source 64-bit RISC-V Core

Fourth Workshop on Computer Architecture Research with RISC-V (CARRV 2020)

May 29<sup>th</sup>, 2020


Nils Wistoff

Moritz Schneider

Frank K. Gürkaynak

Luca Benini

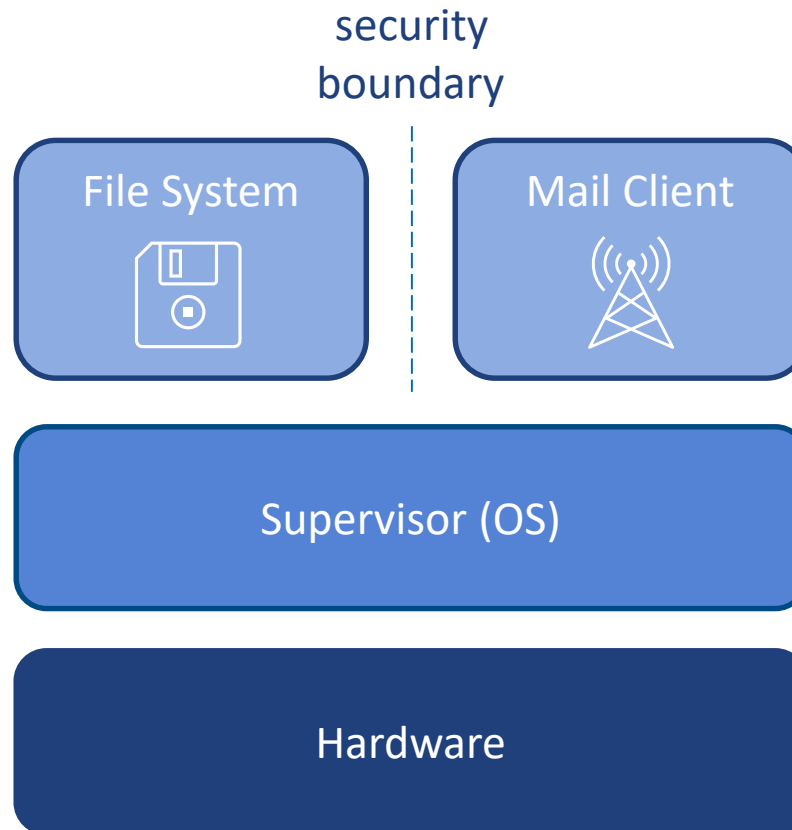
Gernot Heiser



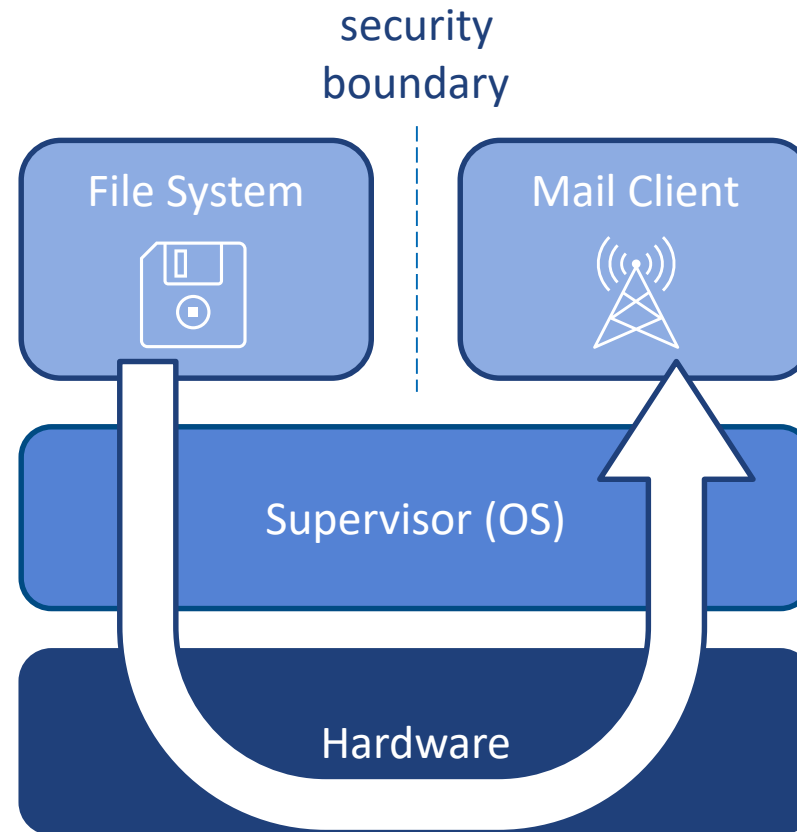
# Outline

1. Covert channels?
2. Measure
3. Mitigate
4. Costs
5. Conclusion

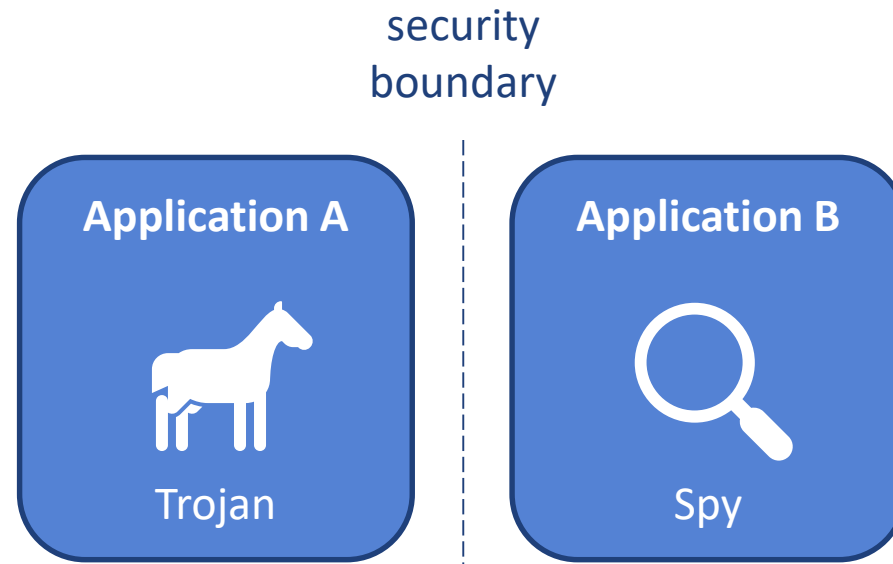
# Covert Channel



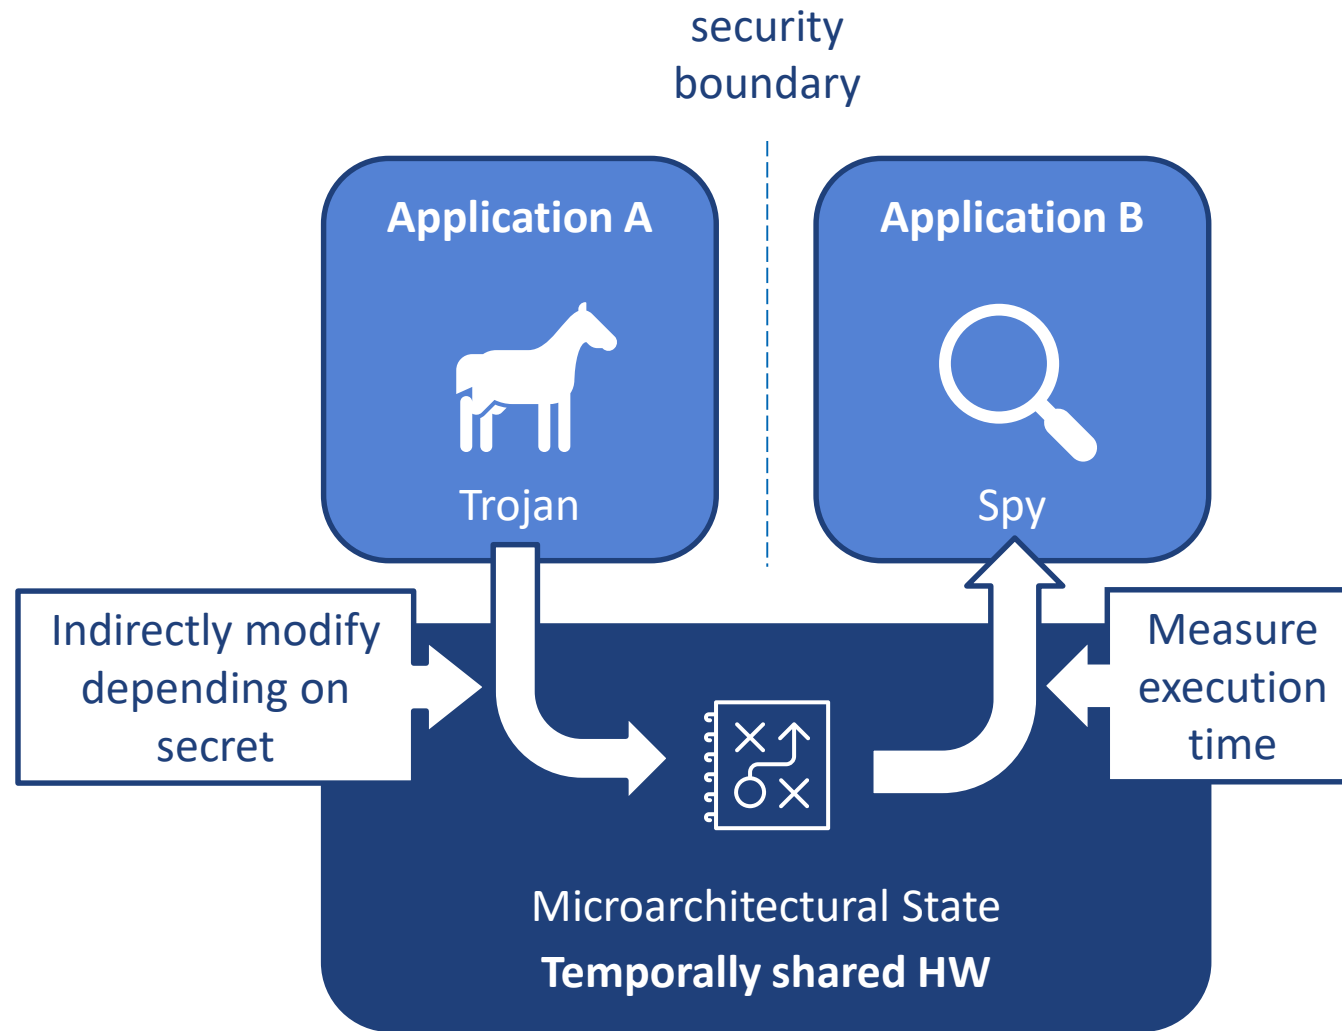
# Covert Channel





# Microarchitectural Timing Channel



# Microarchitectural Timing Channel



# Example: D\$ Timing Channel



(1) Spy:  
Prime


(2) OS:  
Cont. sw.

(3) Trojan:  
Encode s

(4) OS:  
Cont. sw.


(5) Spy:  
Probe

# Example: D\$ Timing Channel – Prime







# Example: D\$ Timing Channel – Prime



# Example: D\$ Timing Channel – Context switch



(1) Spy:  
Prime

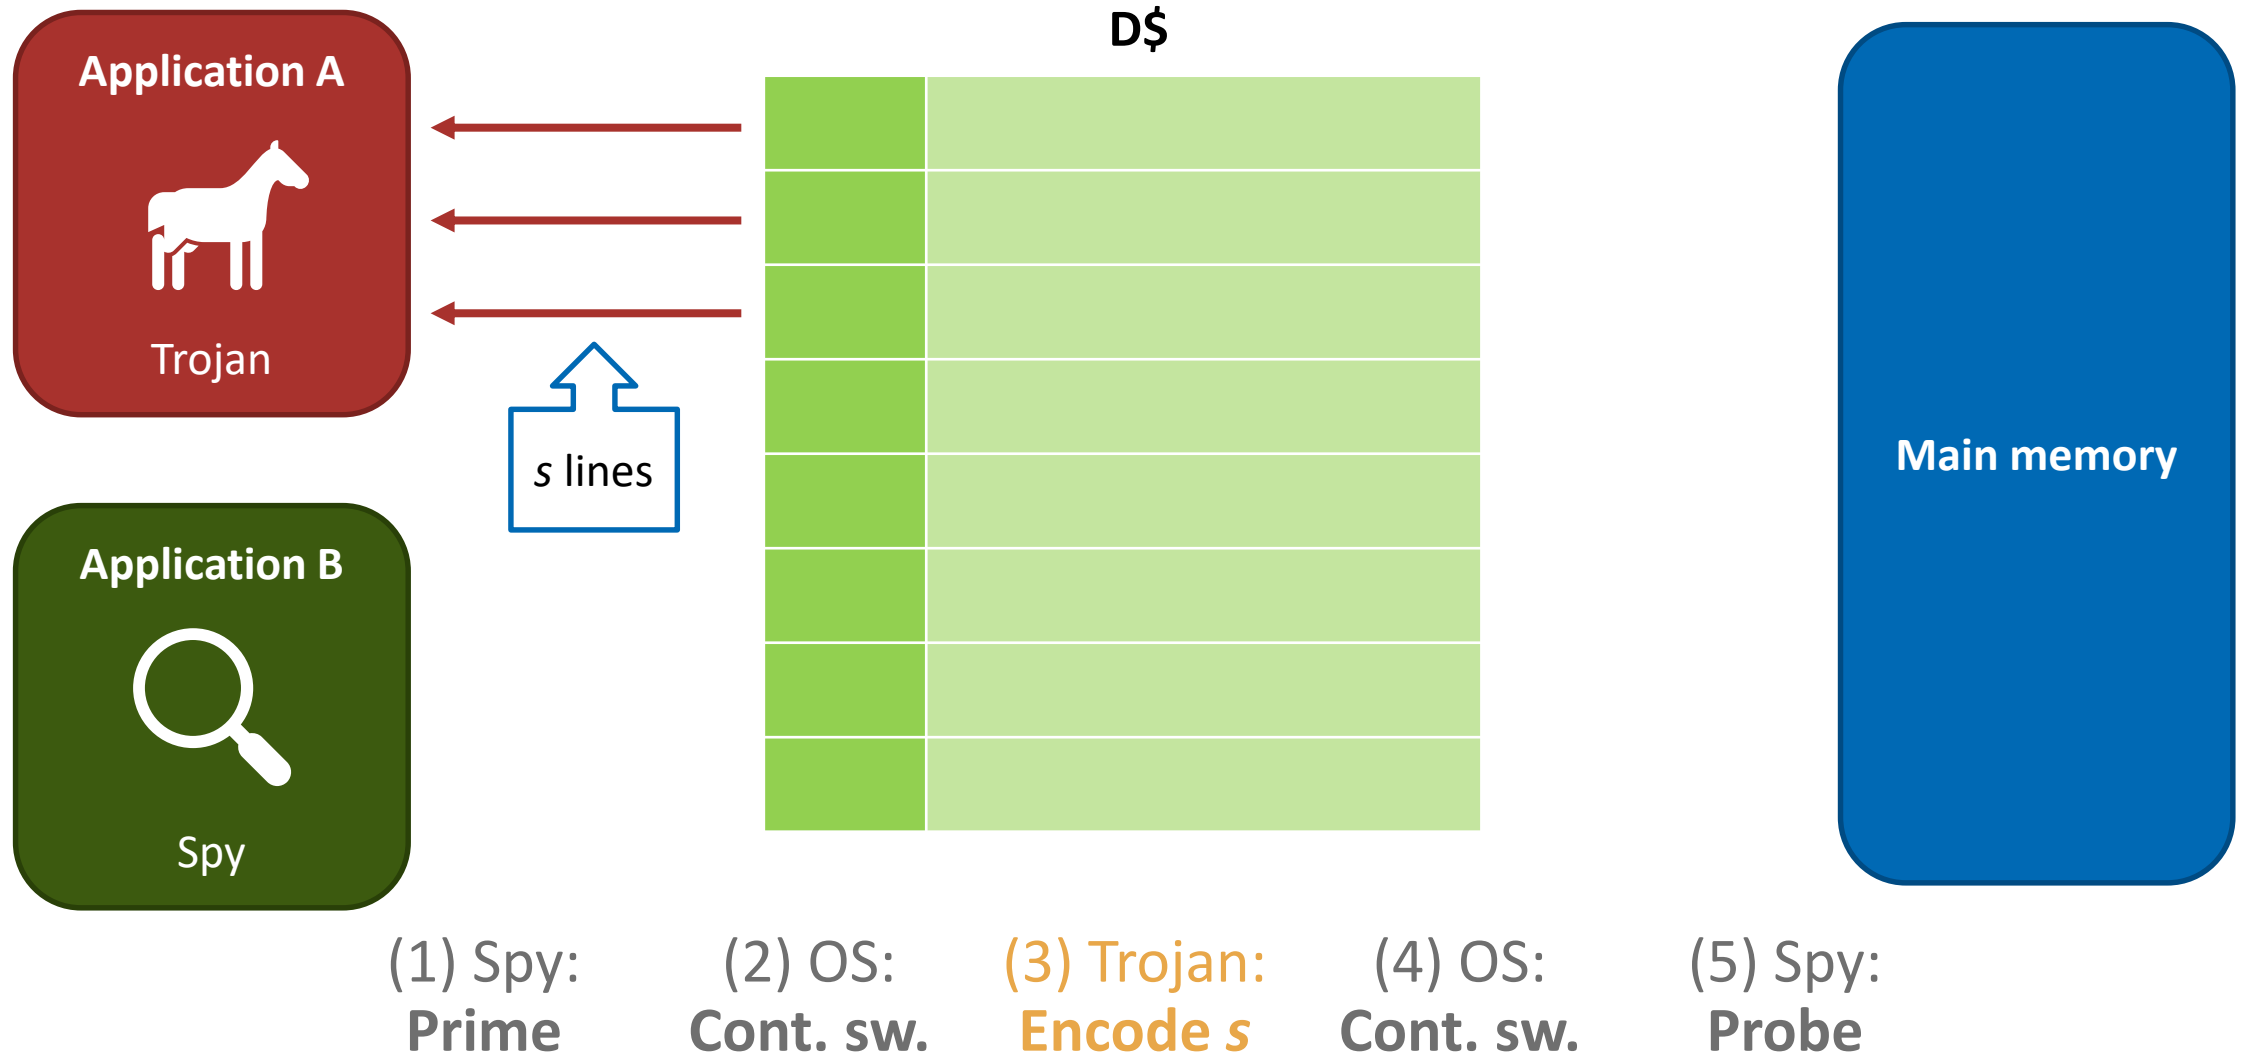
(2) OS:  
Cont. sw.

(3) Trojan:  
Encode s

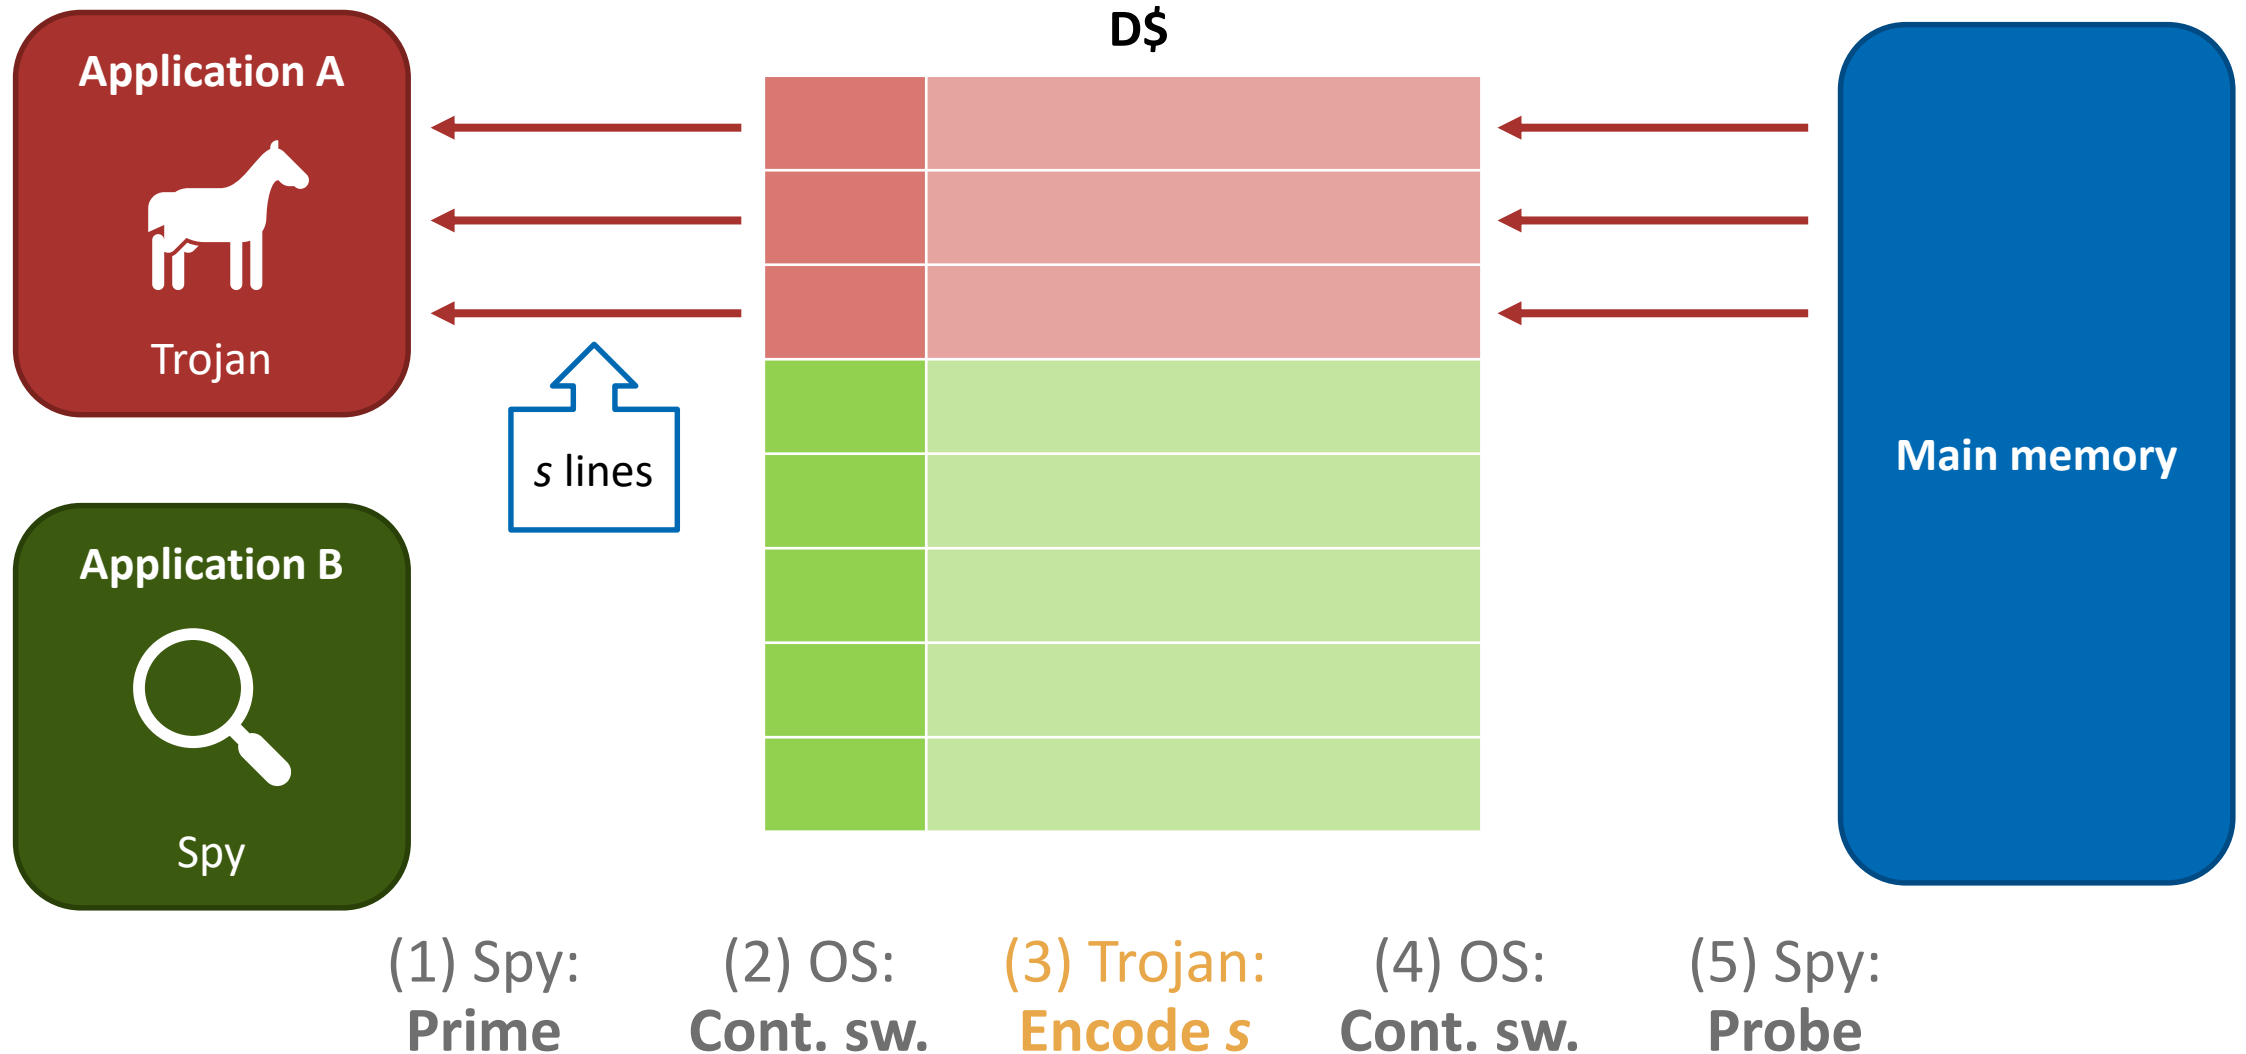
(4) OS:  
Cont. sw.

(5) Spy:  
Probe



# Example: D\$ Timing Channel – Encode $s$



# Example: D\$ Timing Channel – Encode $s$



# Example: D\$ Timing Channel – Context Switch



(1) Spy:  
Prime


(2) OS:  
Cont. sw.

(3) Trojan:  
Encode s

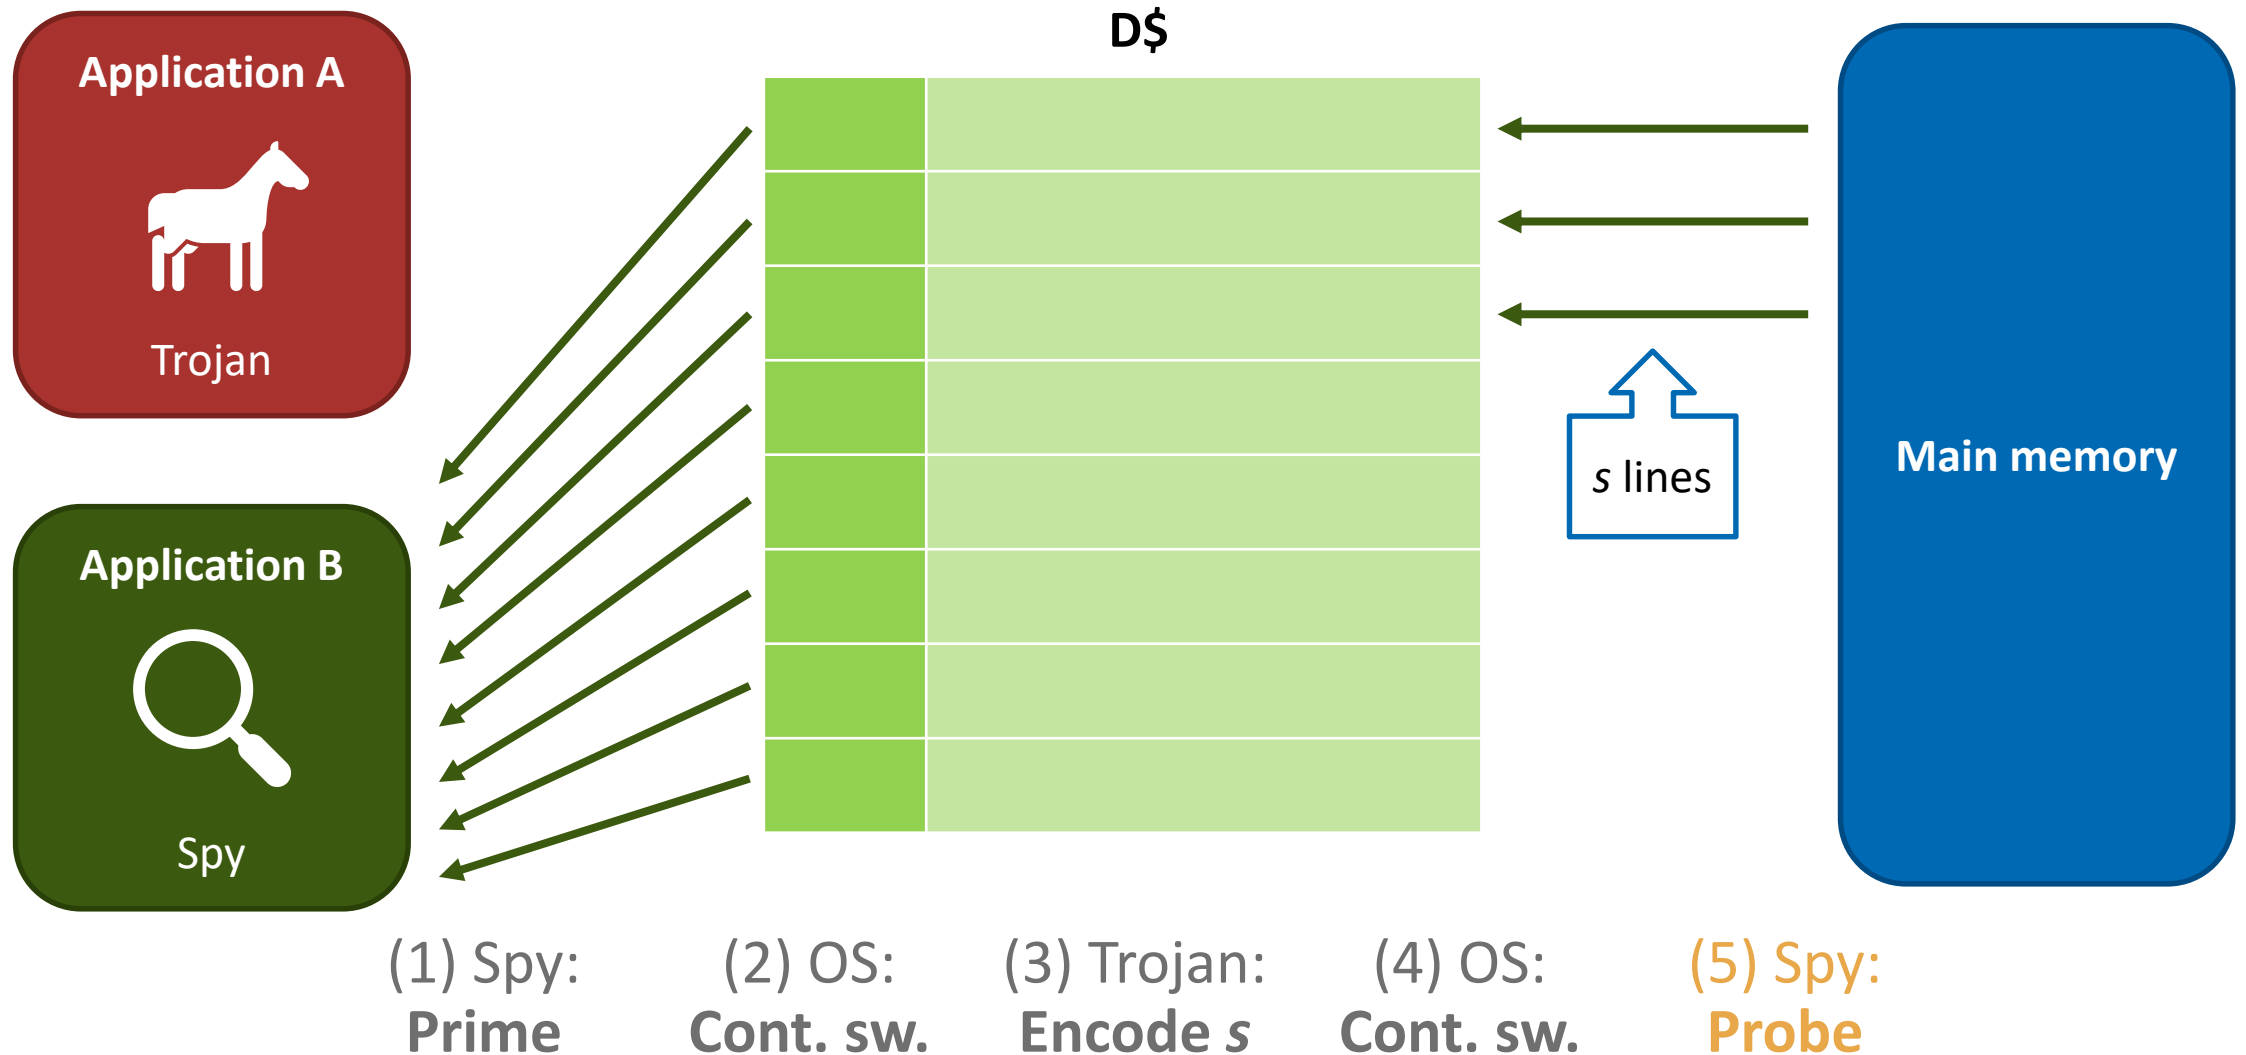
(4) OS:  
Cont. sw.

(5) Spy:  
Probe

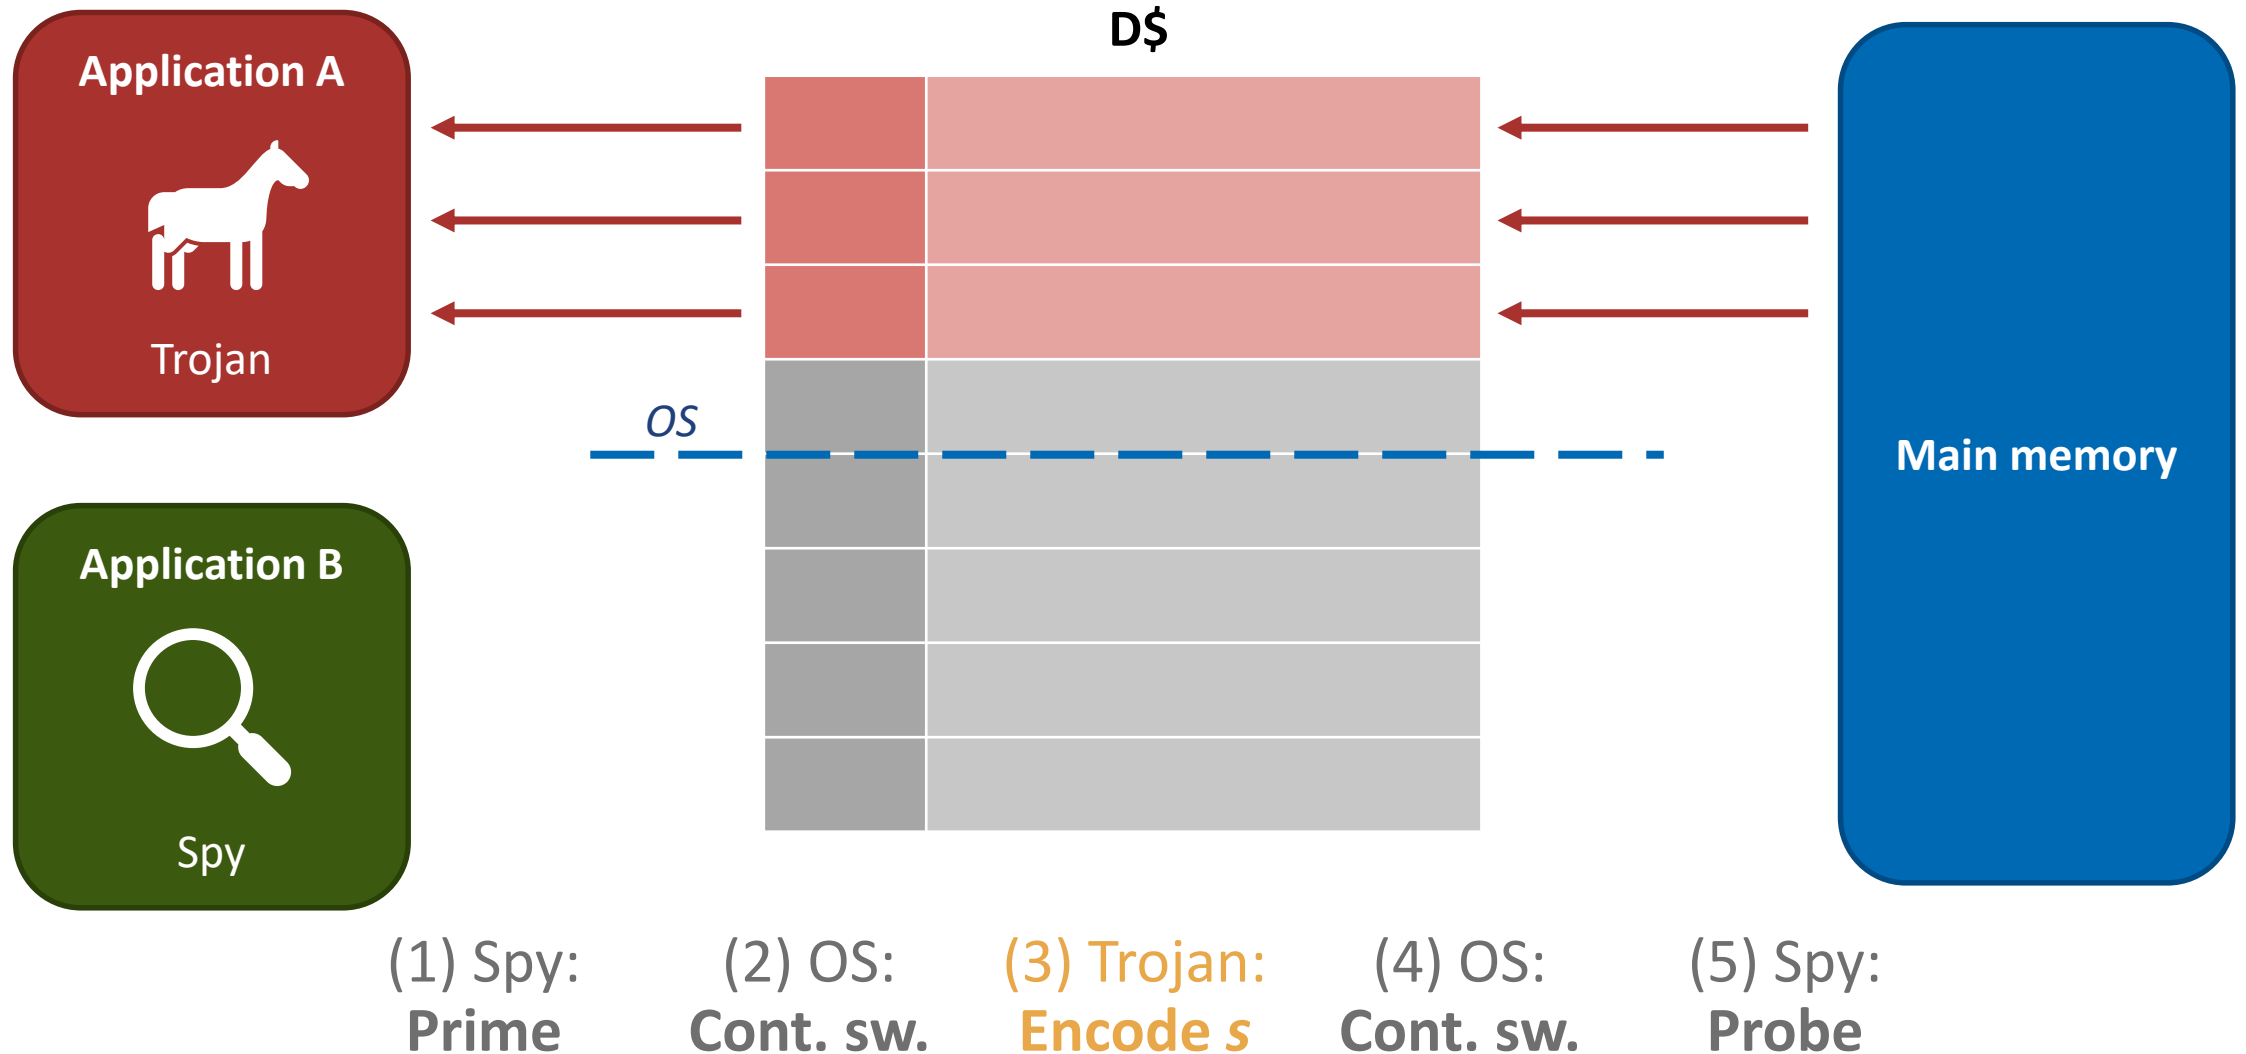
# Example: D\$ Timing Channel – Probe



# Example: D\$ Timing Channel – Probe




# Spatial Partitioning








# Spatial Partitioning



# Temporal Partitioning



# Temporal Partitioning



(1) Spy:  
Prime


(2) OS:  
Cont. sw.

(3) Trojan:  
Encode s



(4) OS:  
Cont. sw.

(5) Spy:  
Probe

# Temporal Partitioning



# Temporal Partitioning



(1) Spy:  
Prime

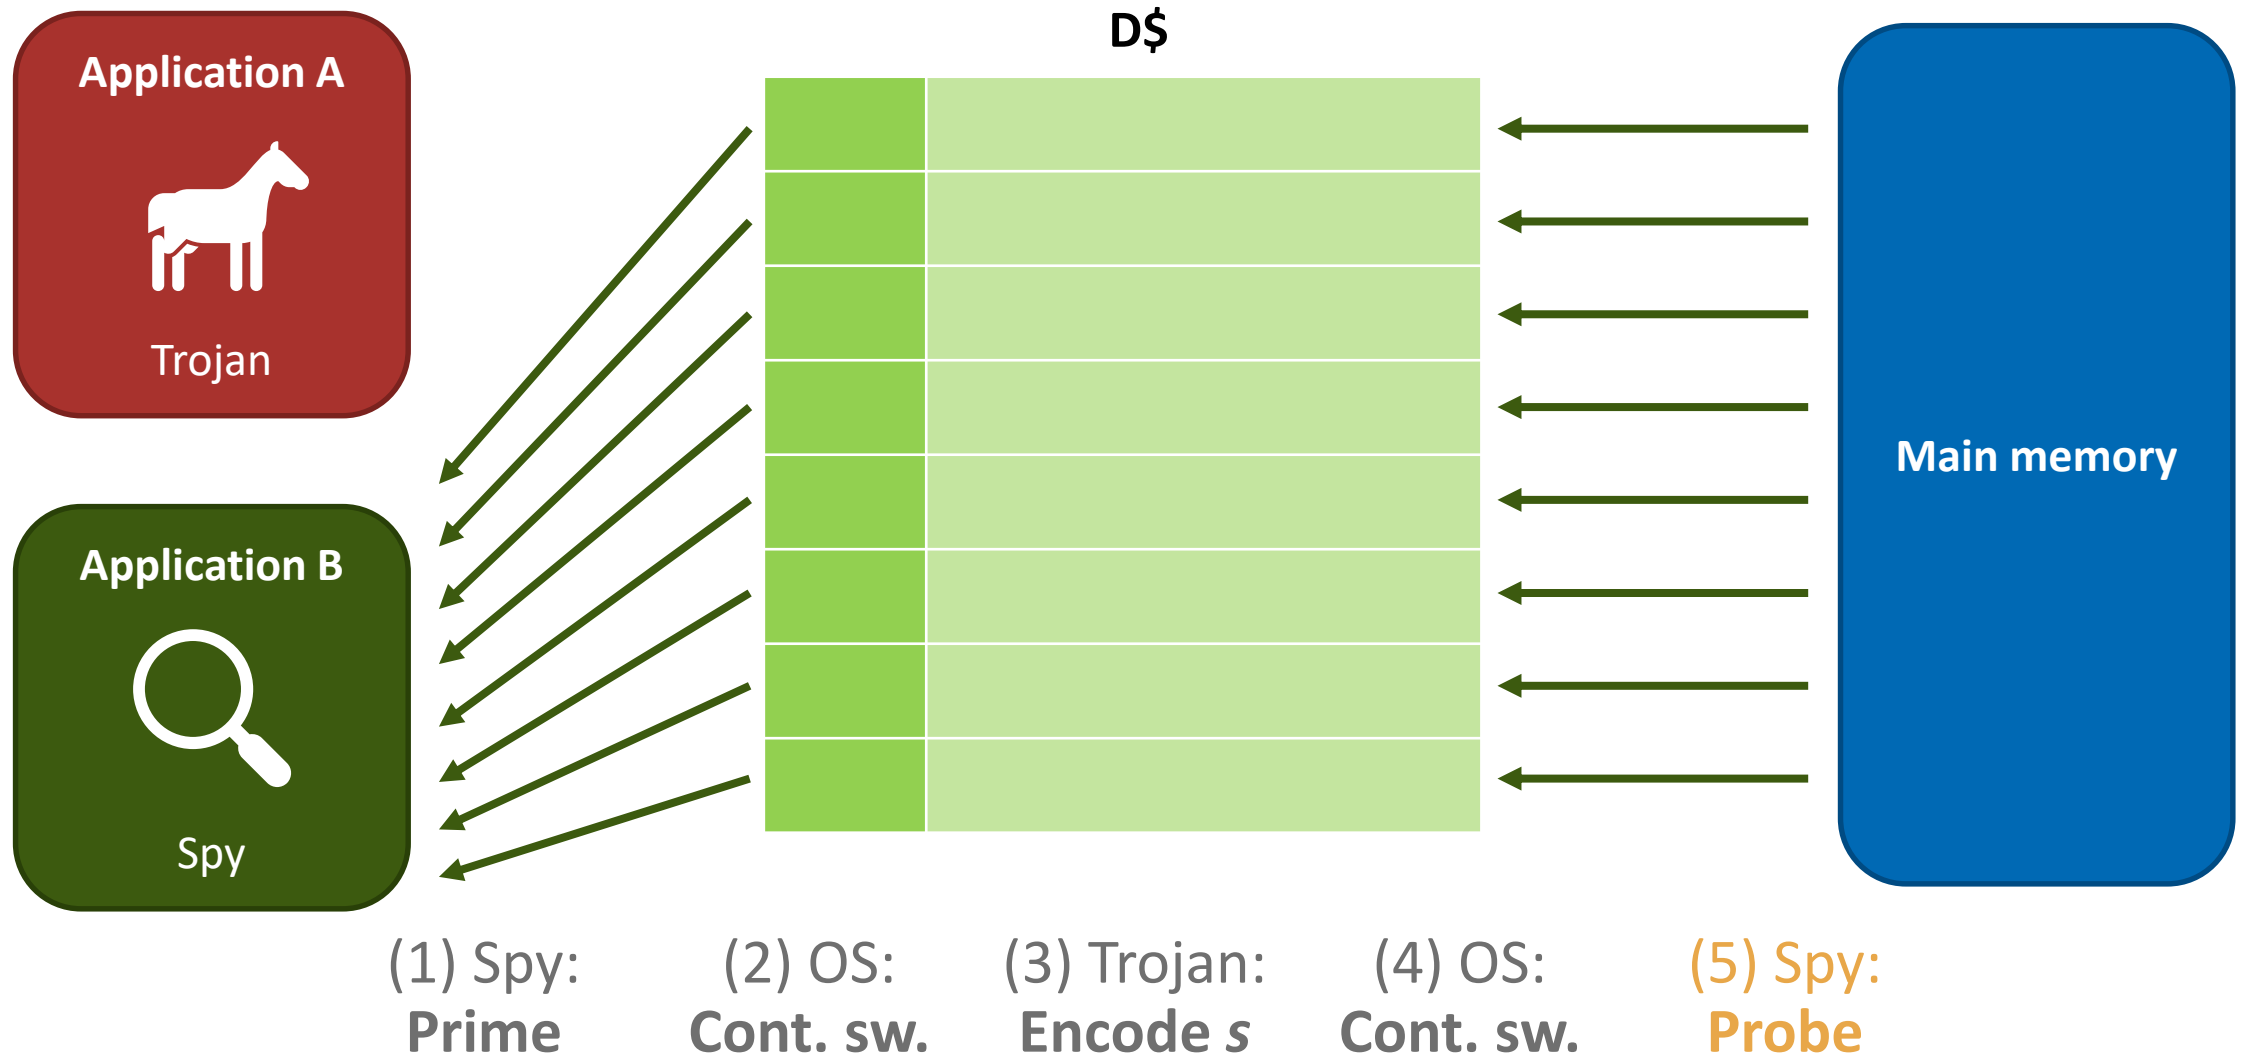
(2) OS:  
Cont. sw.

(3) Trojan:  
Encode s



(4) OS:  
Cont. sw.

(5) Spy:  
Probe

# Temporal Partitioning



# Temporal Partitioning



(1) Spy:  
Prime

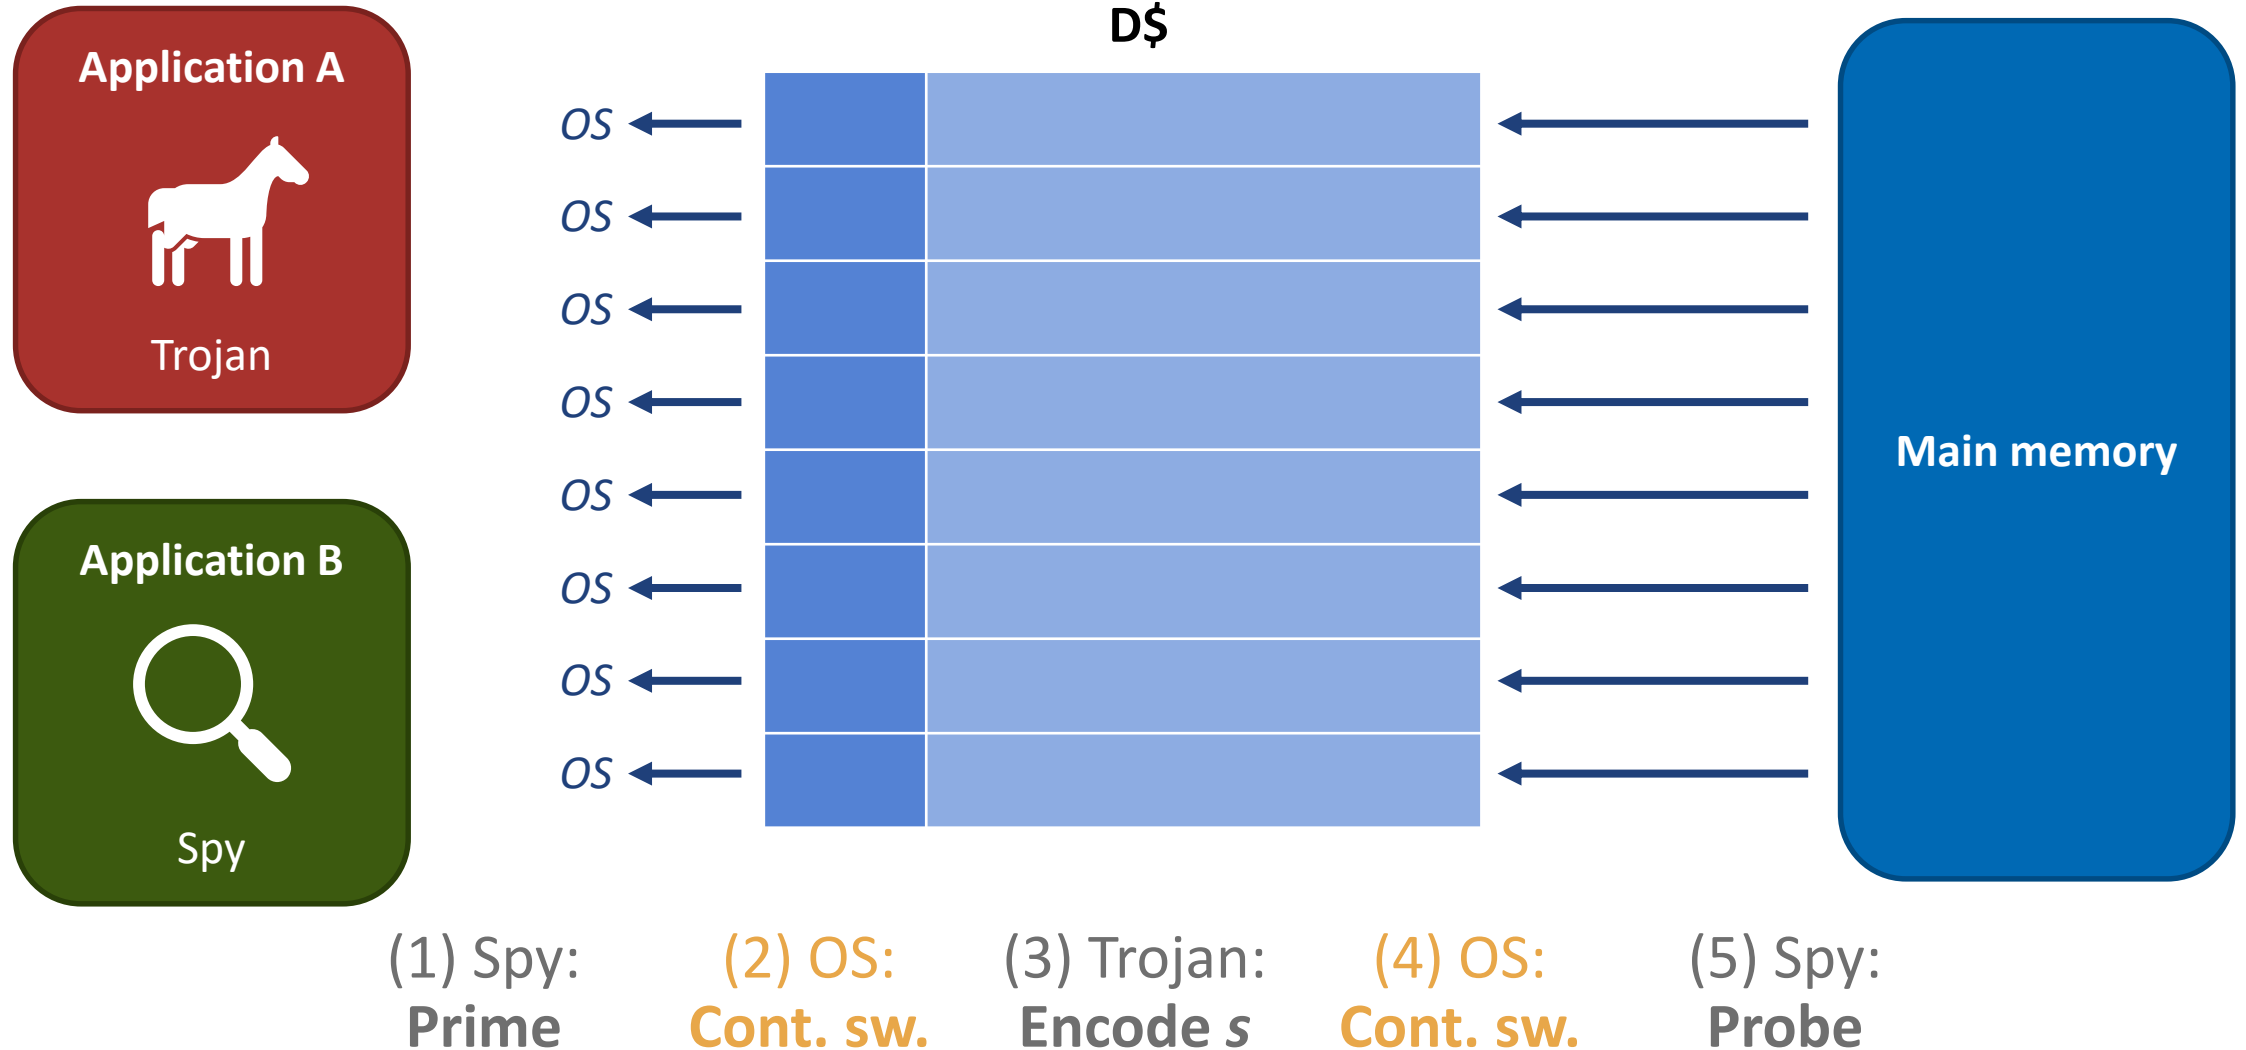
(2) OS:  
Cont. sw.

(3) Trojan:  
Encode s

(4) OS:  
Cont. sw.

(5) Spy:  
Probe

# Flush: SW Approach





# Evaluation Platform

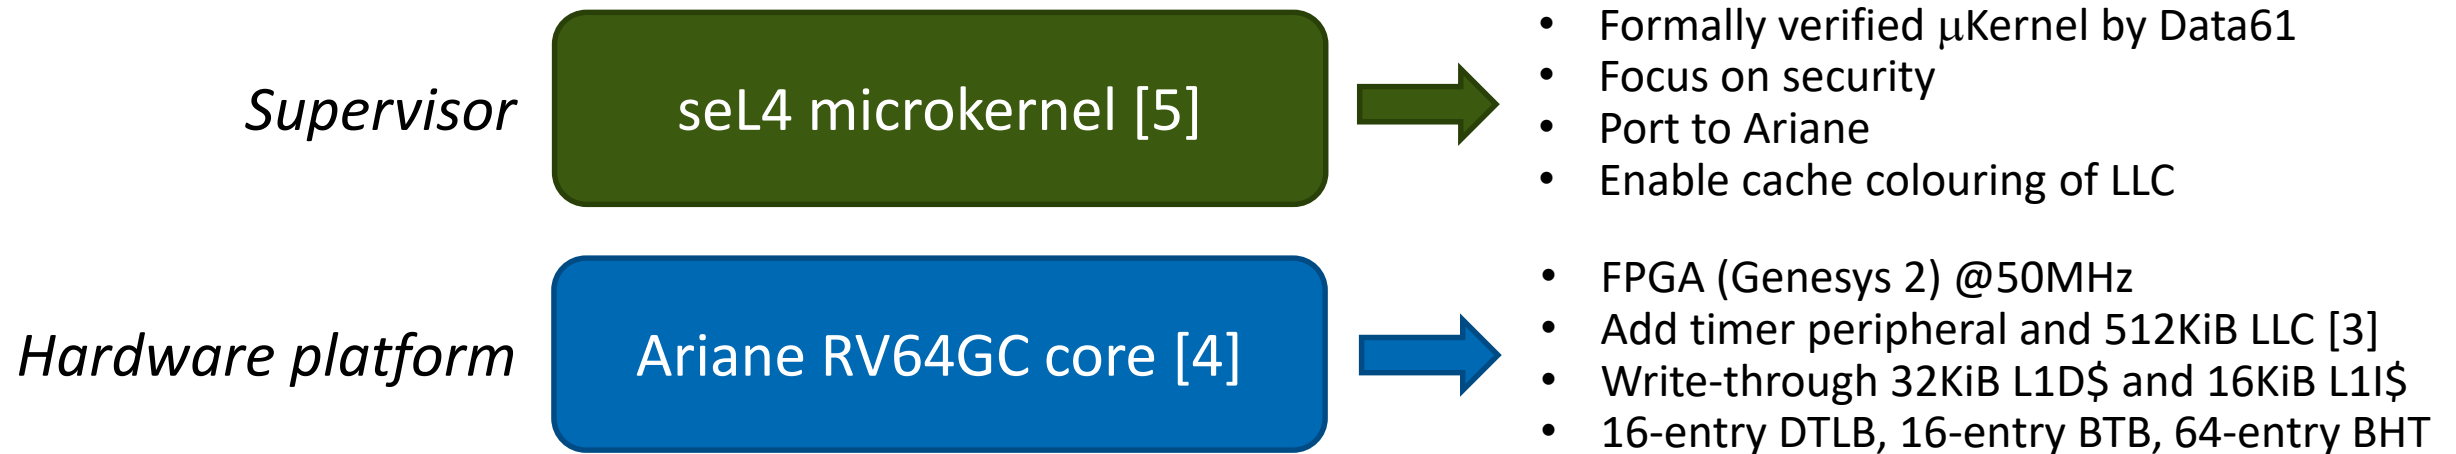
*Hardware platform*

Ariane RV64GC core [4]

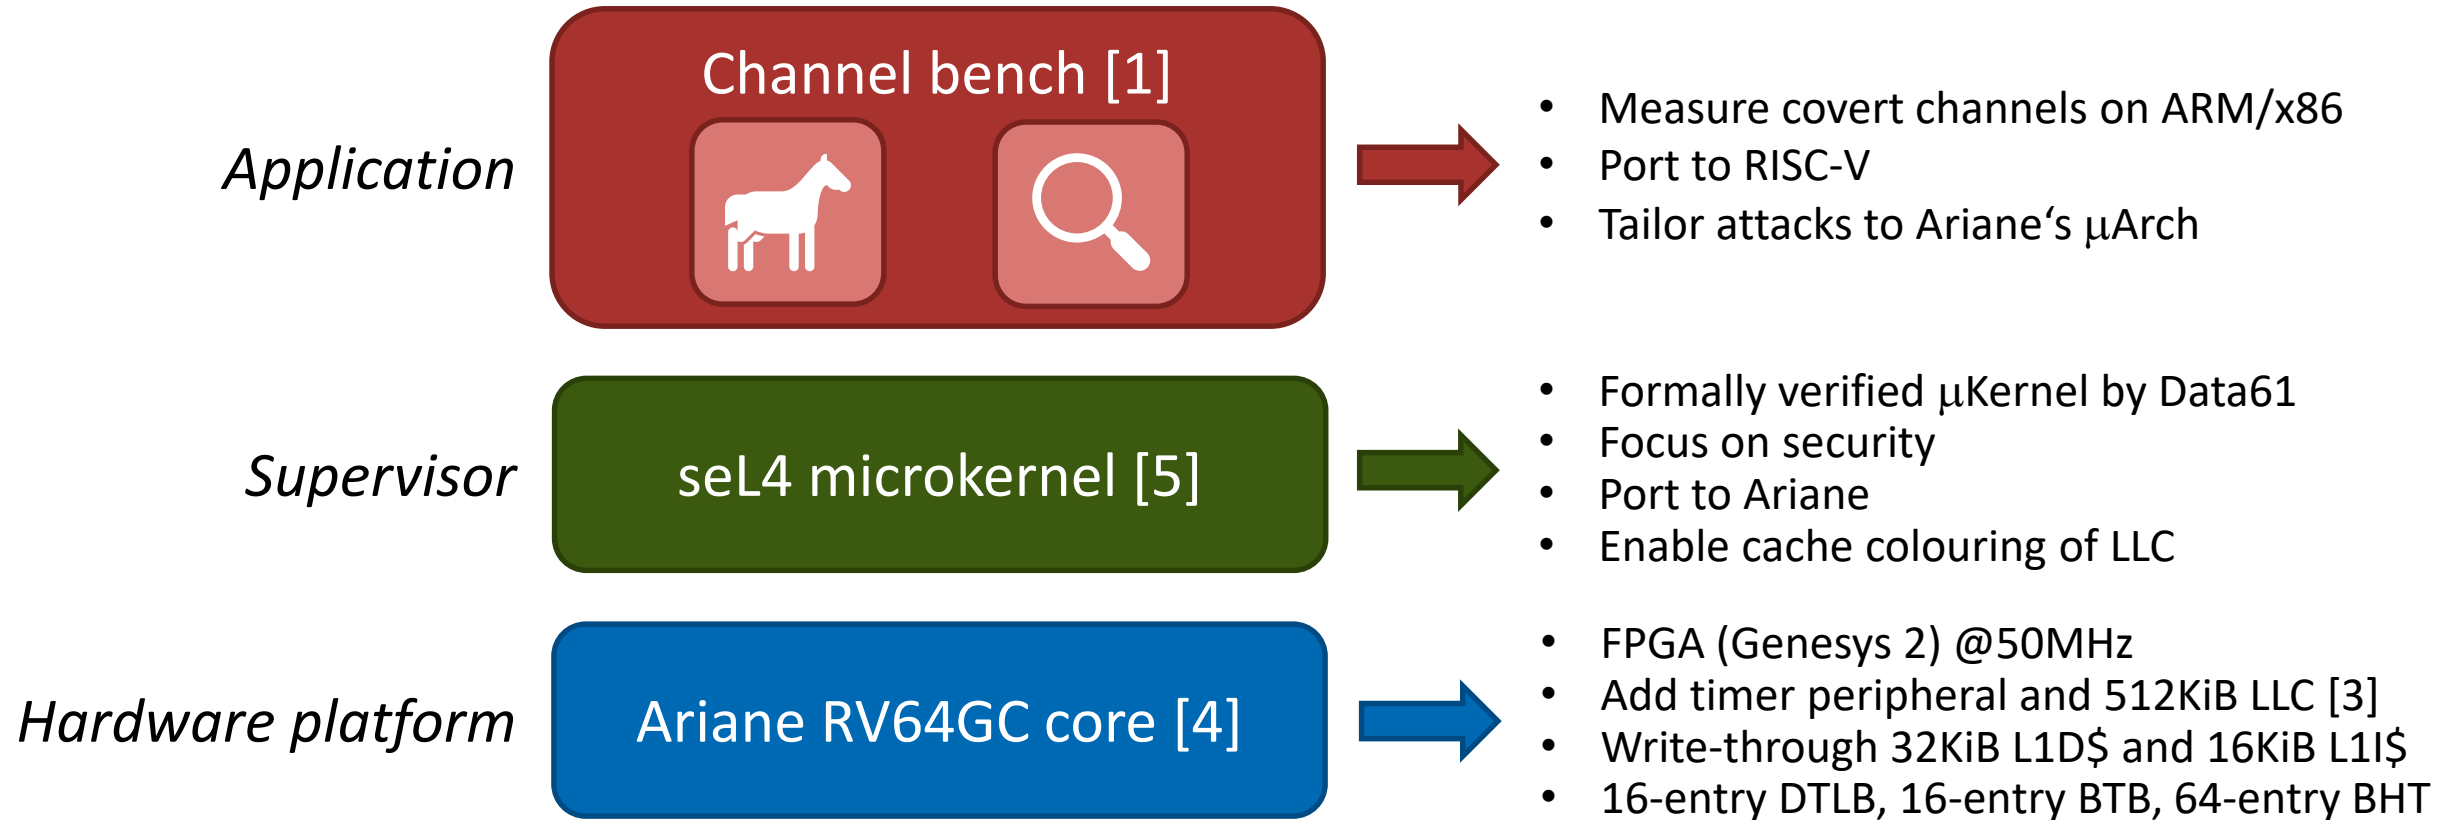


- FPGA (Genesys 2) @50MHz
- Add timer peripheral and 512KiB LLC [3]
- Write-through 32KiB L1D\$ and 16KiB L1I\$
- 16-entry DTLB, 16-entry BTB, 64-entry BHT

# Evaluation Platform



# Evaluation Platform




# Channel Bench Output: L1 D\$

$s_0$	107	$t_0$	83316
$s_1$	11	$t_1$	80209
$s_2$	112	$t_2$	82069
$s_3$	235	$t_3$	88152
$s_4$	246	$t_4$	88856
$s_5$	152	$t_5$	86627
⋮		⋮	


# Channel Matrix: L1 D\$

$N = 10^6$



# Channel Matrix: L1 D\$


$N = 10^6$



# Channel Matrix: L1 D\$

$N = 10^6$

$M = 1667.3$  mb




# Channel Bench Output: L1 D\$

$s_0$	107	$t_0$	83316
$s_1$	11	$t_1$	80209
$s_2$	112	$t_2$	82069
$s_3$	235	$t_3$	88152
$s_4$	246	$t_4$	88856
$s_5$	152	$t_5$	86627


$M$



# Channel Bench Output: L1 D\$



# Channel Bench Output: L1 D\$




$M_0$ : 95% confidence interval of  $M_0^*$   
 $M > M_0 \Rightarrow$  covert channel!

# Channel Matrix: L1 D\$


$$N = 10^6$$

$$M = 1667.3 \text{ mb}$$

$$M_0 = 0.5 \text{ mb}$$




# Flush: SW Approach




# Software Mitigation: L1 D\$ Channel

## Unmitigated



$N = 10^6$ ,  $M = 1667.3$  mb,  $M_0 = 0.5$  mb


## L1 D\$ prime on context switch



$N = 10^6$ ,  $M = 1471.5$  mb,  $M_0 = 0.6$  mb


# Software Mitigation: L1 D\$ Channel

## Single L1 D\$ prime on context switch




$N = 10^6$ ,  $M = 1471.5$  mb,  $M_0 = 0.6$  mb

## Double L1 D\$ prime on context switch

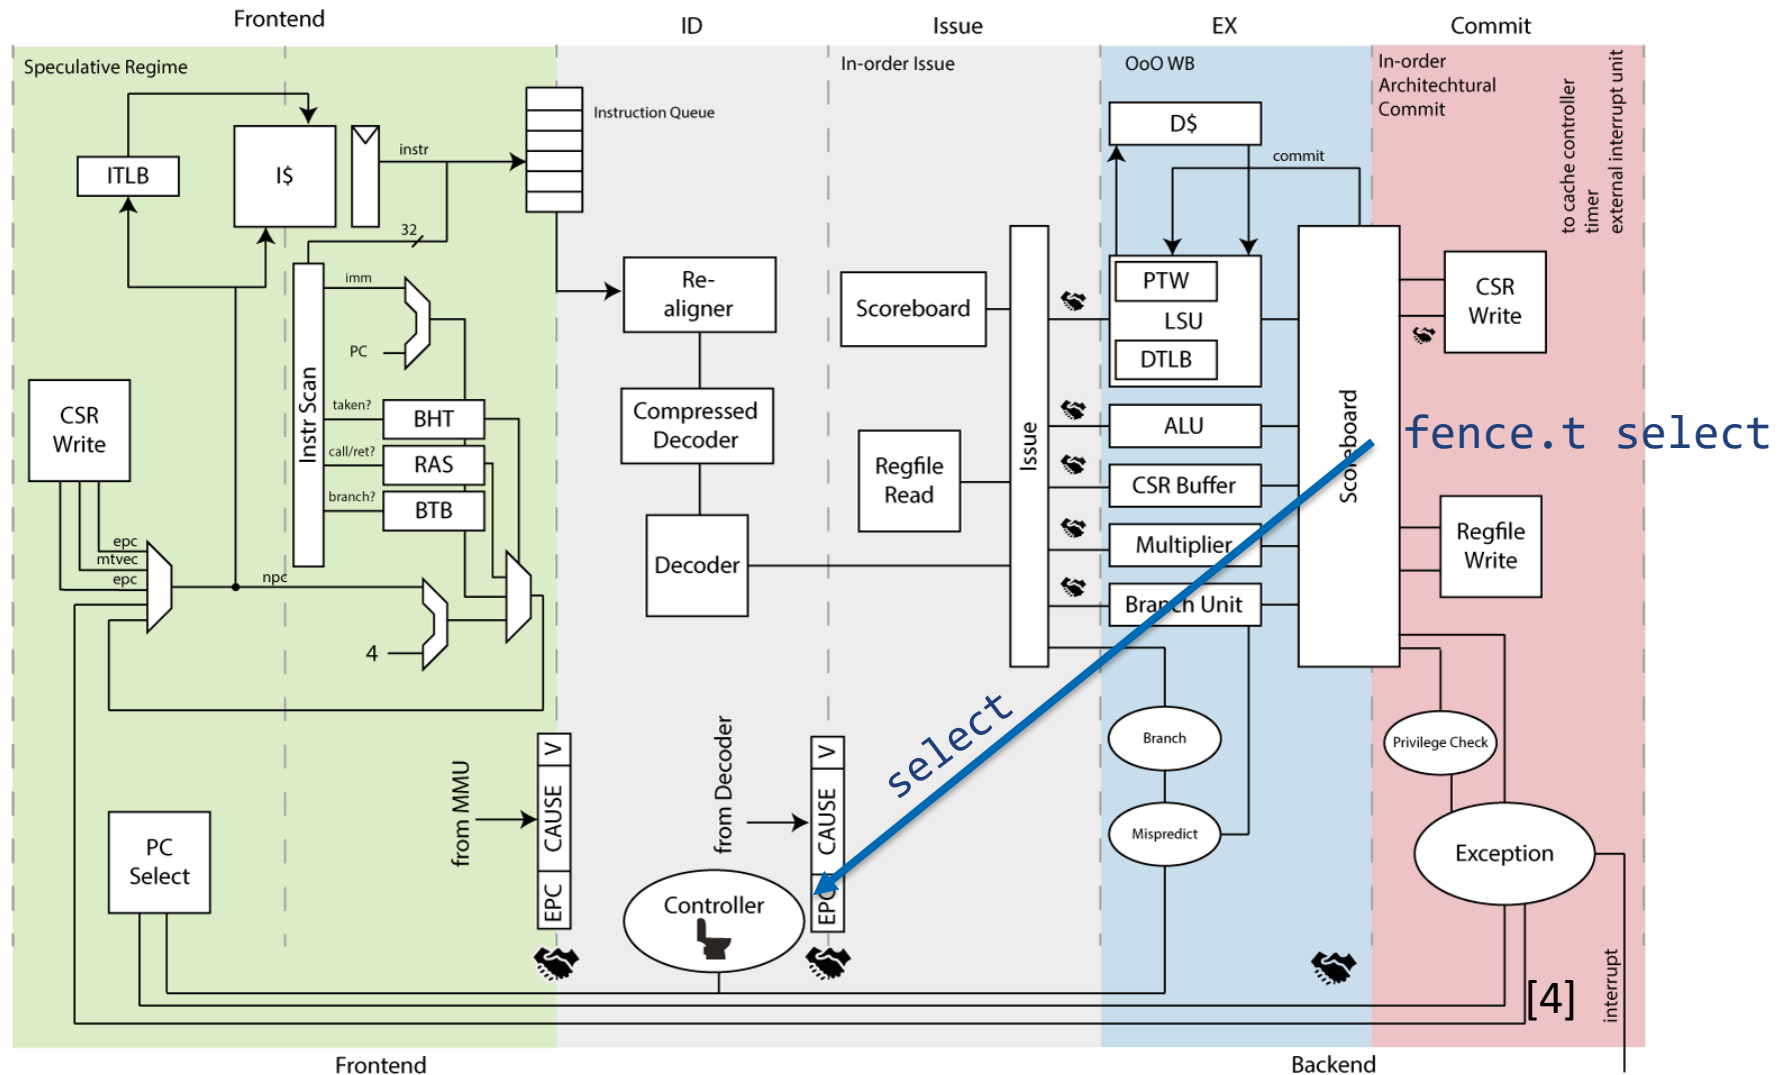


$N = 10^6$ ,  $M = 515.7$  mb,  $M_0 = 1.1$  mb

# Temporal Fence Instruction (`fence.t`)

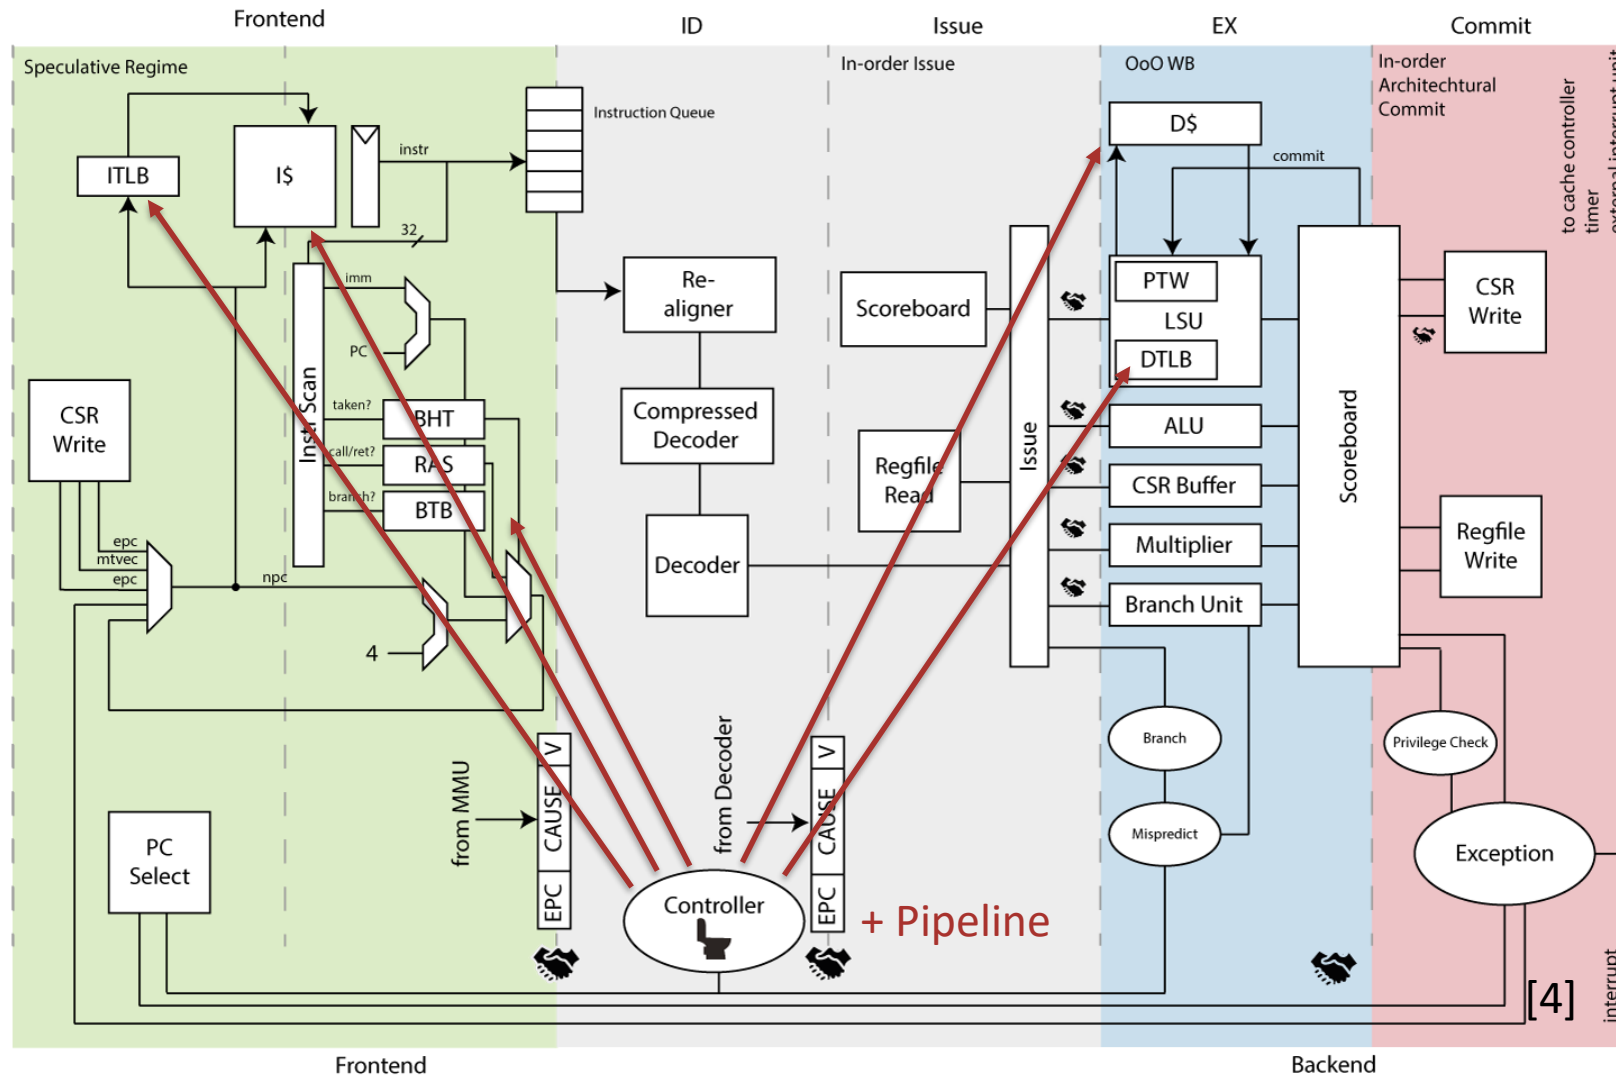


# Temporal Fence Instruction (fence.t)






# Temporal Fence Instruction (fence.t)




# fence.t: L1 D\$ Channel

## Unmitigated



$N = 10^6$ ,  $M = 1667.3$  mb,  $M_0 = 0.5$  mb


## Flush targeted components on context switch



$N = 10^6$ ,  $M = 7.7$  mb,  $M_0 = 1.4$  mb


# fence.t: L1 D\$ Channel

## Unmitigated



$N = 10^6$ ,  $M = 1667.3$  mb,  $M_0 = 0.5$  mb

## Flush targeted components on context switch



... but wait!


$N = 10^6$ ,  $M = 7.7$  mb,  $M_0 = 1.4$  mb

# Vulnerable 2<sup>nd</sup> Order State-Holding Components

- **L1 D\$:**
  - LFSR for pseudo-random replacement policy
  - Memory arbiter
  - TX FIFO
  - Write-buffer arbiters
- **L1 I\$:**
  - LFSR for pseudo-random replacement policy
- **TLBs:**
  - Pseudo-LRU tree for replacement policy


# Full fence.t: L1 D\$ Channel

Unmitigated



$N = 10^6$ ,  $M = 1667.3$  mb,  $M_0 = 0.5$  mb


Flush all vulnerable components  
on context switch



$N = 10^6$ ,  $M = 8.4$  mb,  $M_0 = 9.6$  mb


# L1 I\$ Channel

## Unmitigated



$N = 10^6$ ,  $M = 1905.0$  mb,  $M_0 = 0.5$  mb


## Flush all vulnerable components on context switch



$N = 10^6$ ,  $M = 19.5$  mb,  $M_0 = 20.5$  mb


# TLB Channel

## Unmitigated



$N = 10^6$ ,  $M = 409.2$  mb,  $M_0 = 0.1$  mb


## Flush all vulnerable components on context switch



$N = 10^6$ ,  $M = 2.7$  mb,  $M_0 = 5.4$  mb


# BTB Channel

## Unmitigated



$N = 10^6$ ,  $M = 3481.3$  mb,  $M_0 = 0.1$  mb

## Flush all vulnerable components on context switch




$N = 10^6$ ,  $M = 33.0$  mb,  $M_0 = 57.6$  mb




# BHT Channel

## Unmitigated



$N = 10^6$ ,  $M = 4873.3$  mb,  $M_0 = 0.1$  mb

## Flush all vulnerable components on context switch



$N = 10^6$ ,  $M = 44.1$  mb,  $M_0 = 58.8$  mb

# Context Switch Latency

seL4 one-way inter-address-space IPC microbenchmark

Unmitigated	
Hot	Cold
430 ( $\pm 7.0$ )	1,180 ( $\pm 1.0$ )

# Context Switch Latency

## seL4 one-way inter-address-space IPC microbenchmark

Unmitigated		D\$ Software Flush	
Hot	Cold	Single	Double
430 ( $\pm 7.0$ )	1,180 ( $\pm 1.0$ )	12,099 ( $\pm 52$ )	51,876 ( $\pm 256$ )

# Context Switch Latency

## seL4 one-way inter-address-space IPC microbenchmark

Unmitigated		D\$ Software Flush		HW Flush
Hot	Cold	Single	Double	
430 ( $\pm 7.0$ )	1,180 ( $\pm 1.0$ )	12,099 ( $\pm 52$ )	51,876 ( $\pm 256$ )	1,502 ( $\pm 0.9$ )

# Hardware Costs: FPGA

	LUTs	Registers	Muxes
Unmodified	102,796 ( $\pm 10$ )	58,957 ( $\pm 208$ )	13,590 ( $\pm 38$ )
w/ fence.t	102,792 ( $\pm 57$ )	60,607 ( $\pm 5$ )	15,038 ( $\pm 2$ )
	$\pm 0\%$	+2.8%	+10.6%

# Conclusion

- **We measure five distinct covert channels on Ariane**
- **Confirmed: OS needs HW-support for time protection [1]**
- **HW-mechanism must flush *all*  $\mu$ Arch state**
  - Identifying  $\mu$ Arch state not always straight-forward
  - Systematic approach for HW / Security codesign needed
- **Further, off-core covert channels still need to be addressed**
  - e.g. DRAM, thermal controller, etc.

# Sources

- [1] Qian Ge, Yuval Yarom, Tom Chothia, and Gernot Heiser: “Time Protection: The Missing OS Abstraction”, EuroSys, 2019
- [2] R. E. Kessler and Mark D. Hill: “Page Placement Algorithm for Large Real-Indexed Caches”, ACM Trans. Comp. Syst. 19, 1992
- [3] Wolfgang Rönninger: “Memory Subsystem for the First Fully Open-Source RISC-V Heterogeneous SoC”, Master’s thesis, ETH Zurich, 2019
- [4] Florian Zaruba and Luca Benini: “The Cost of Application-Class Processing: Energy and Performance Analysis of a Linux-Ready 1.7-GHz 64-Bit RISC-V Core in 22-nm FDSOI Technology”, IEEE Trans. on VLSI Systems 27, 2019
- [5] Gerwin Klein, June Andronick, Kevin Elphistone, Toby Murray, Thomas Sewell, Rafal Kolanski, and Gernot Heiser: “Comprehensive Formal Verification of an OS Microkernel”, ACM Trans. Comp. Syst. 32, 2014

# Prevention of Microarchitectural Covert Channels on an Open-Source 64-bit RISC-V Core

Fourth Workshop on Computer Architecture Research with RISC-V (CARRV 2020)

May 29<sup>th</sup>, 2020


Nils Wistoff

Moritz Schneider

Frank K. Gürkaynak

Luca Benini

Gernot Heiser

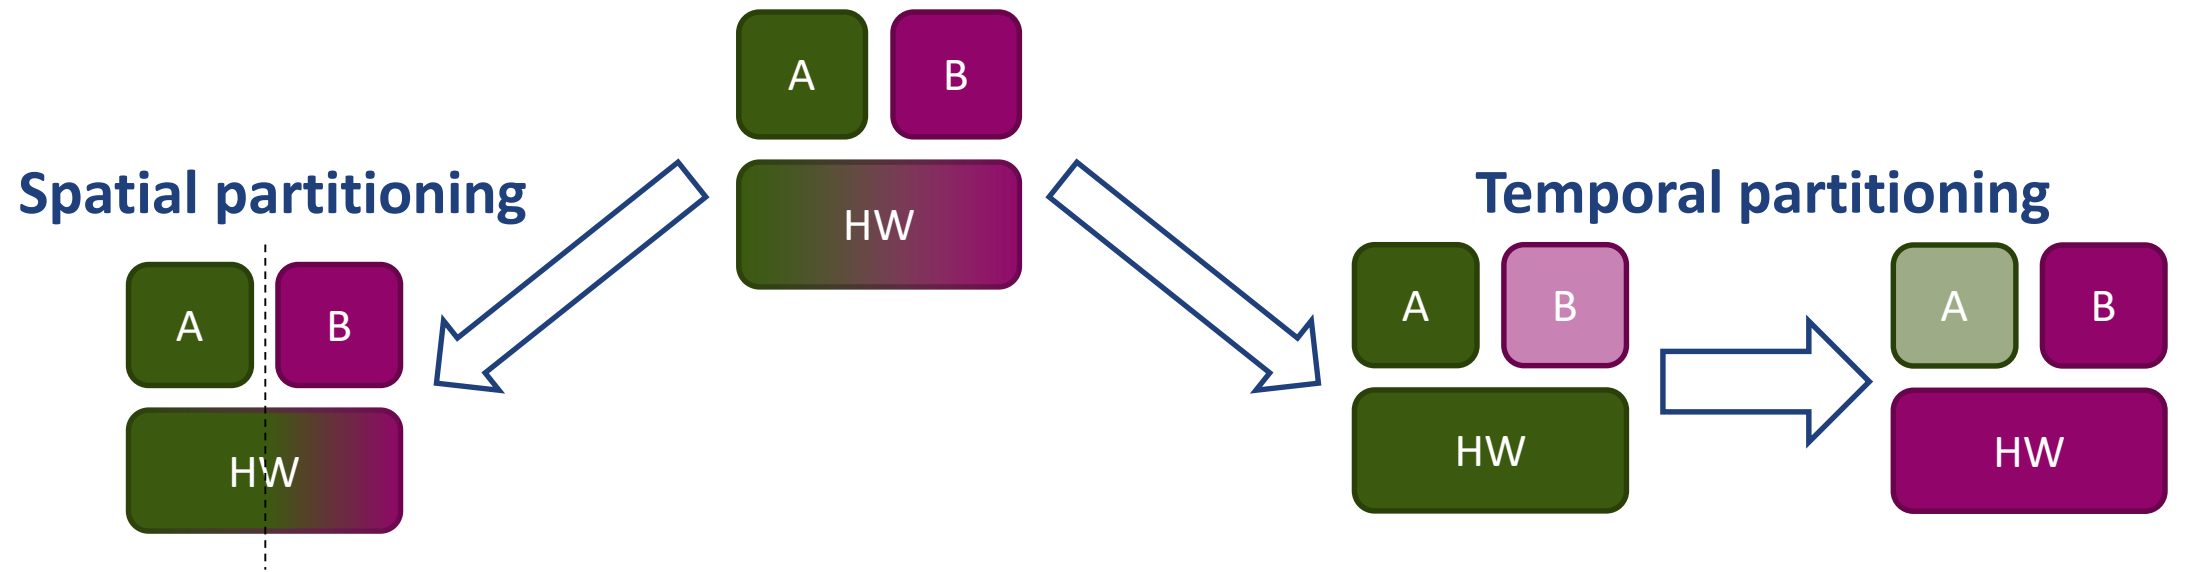




# Hardware Costs: FPGA

	LUTs	Registers	Muxes
Unmodified	102,796 ( $\pm 10$ )	58,957 ( $\pm 208$ )	13,590 ( $\pm 38$ )
w/ fence.t	102,792 ( $\pm 57$ ) 50.4%	60,607 ( $\pm 5$ ) 14.9%	15,038 ( $\pm 2$ ) 9.8%
	$\pm 0\%$	+2.8%	+10.6%

# Time Protection [1]



- Off-core components
- e.g. cache colouring (LLC) [2]
- **Not a solution for on-core components!**

- On-core components
- e.g. L1 caches, TLBs, branch predictors
- Reset  $\mu$ Arch state on context switch