

SSE-RV: Secure Speculative Execution via RISC-V Open Hardware Design

Majid Sabbagh, Yunsi Fei, David Kaeli

Department of Electrical and Computer Engineering
Northeastern University, Boston, MA, USA

Fifth Workshop on Computer Architecture Research with RISC-V (CARRV 2021)



Northeastern University



A sample uarch: Intel Skylake

Closed/propriety hardware design

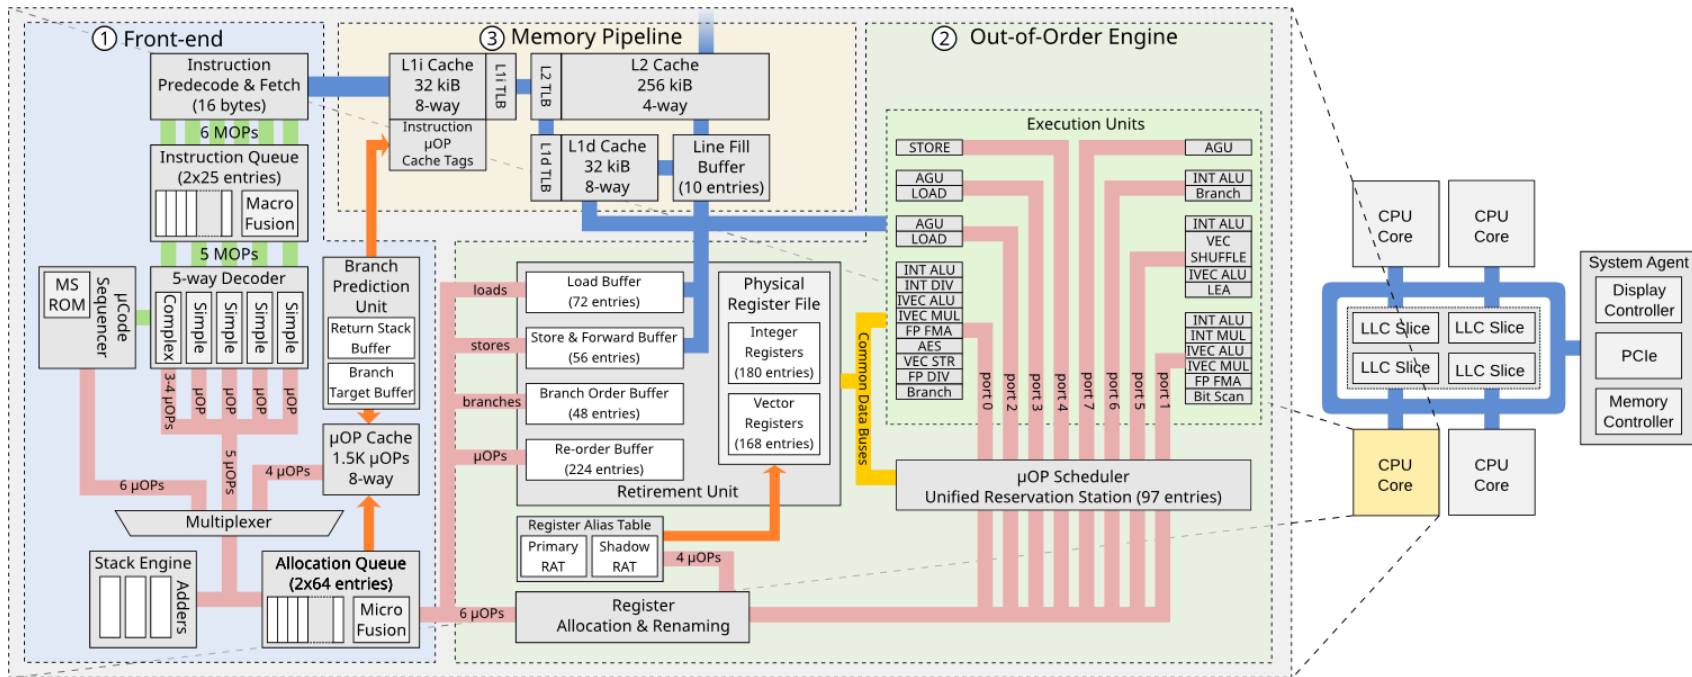


Figure by Stephan van Schaik, *the VU*

Transient-Execution attacks: why do they matter?



Bounds Check Bypass (Spectre v1) CVE-2017-5753 INTEL-SA-00088	Branch Target Injection (Spectre v2) CVE-2017-5715 INTEL-SA-00088	Rogue Data Cache Load (Meltdown) CVE-2017-5754 INTEL-SA-00088	Rogue System Register Read (v3a)-RDMSR CVE-2018-3640 INTEL-SA-00115	Speculative Store Bypass (Spectre v4) CVE-2018-3639 INTEL-SA-00115	L1 Terminal Fault (L1TF) CVE-2018-3615 CVE-2018-3620 CVE-2018-3646 INTEL-SA-00161
Intel Guidance: Bounds Check Bypass Overview	Intel Guidance: Branch Target Injection Overview	Intel Guidance: Rogue Data Cache Load Overview	Intel Guidance: Rogue System Register Read Overview	Intel Guidance: Speculative Store Bypass Overview	Intel Guidance: L1 Terminal Fault Overview

	Link
Intel	Security Advisory / Newsroom / Whitepaper
ARM	Security Update
AMD	Security Information
RISC-V	Blog
NVIDIA	Security Bulletin / Product Security
Microsoft	Security Guidance / Information regarding anti-virus software / Azure Blog / Windows (Client) / Windows (Server)
Amazon	Security Bulletin
Google	Project Zero Blog / Need to know
Android	Security Bulletin
Apple	Apple Support
Lenovo	Security Advisory
IBM	Blog
Dell	Knowledge Base / Knowledge Base (Server)
Hewlett Packard Enterprise	Vulnerability Alert
HP Inc.	Security Bulletin
Huawei	Security Notice
Synology	Security Advisory
Cisco	Security Advisory
FS	Security Advisory
Mozilla	Security Blog
Red Hat	Vulnerability Response / Performance Impacts
Debian	Security Tracker
Ubuntu	Knowledge Base
SUSE	Vulnerability Response
Fedora	Kernel update
Qubes	Announcement
Fortinet	Advisory
NetApp	Advisory
LLVM	Spectre (Variant #2) Patch / Review __builtin_load_no_speculate / Review llvm.nospeculateload
CERT	Vulnerability Note
MITRE	CVE-2017-5715 / CVE-2017-5753 / CVE-2017-5754
VMWare	Security Advisory / Blog
Citrix	Security Bulletin / Security Bulletin (XenServer)
Xen	Security Advisory (XSA-254) / FAQ

ID Eviction Sampling CVE-2020-0543 TEL-SA-00329	Load Value Injection: Stale Data CVE-2020-0551 INTEL-SA-00334	Load Value Injection: Zero Data CVE-2020-0551 INTEL-SA-00334	Snoop-assisted L1D Sampling CVE-2020-0550 INTEL-SA-00330	Special Register Buffer Data Sampling CVE-2020-0543 INTEL-SA-00320	Fast Forward Store Predictor CVE-2020-0543 INTEL-SA-00331
Intel Guidance: ID Eviction Sampling Overview	Intel Guidance: Load Value Injection Deep Dive	Intel Guidance: Load Value Injection Deep Dive	Intel Guidance: Snoop-assisted L1D Sampling Deep Dive	Intel Guidance: SRBDS Deep Dive	Intel Guidance: Updating Microcode

Highlights

❑ Implemented a novel taint tracking architecture

- ❖ Based on *SonicBOOM*

- ❖ **SSE-RV protects against Spectre attacks**, outperforming the state-of-the-art

❑ Developed an **FPGA prototype for the proposed SSE-RV**

❑ **Evaluated the security guarantees** delivered by our protection scheme, as well as **area/power/performance overheads**

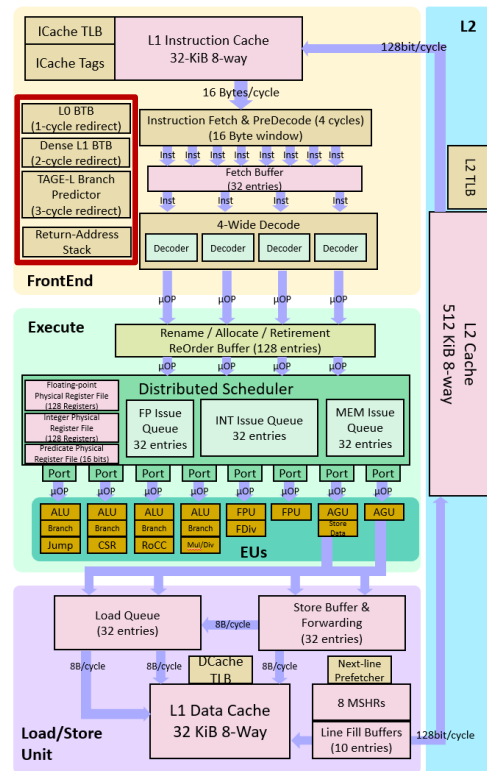
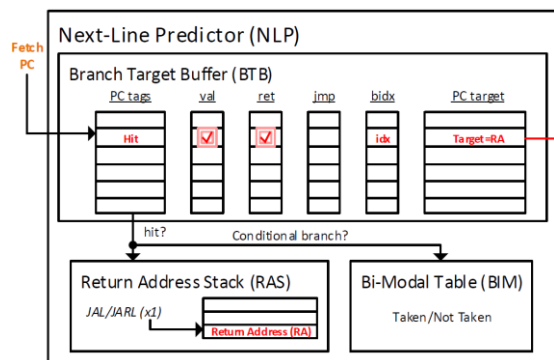
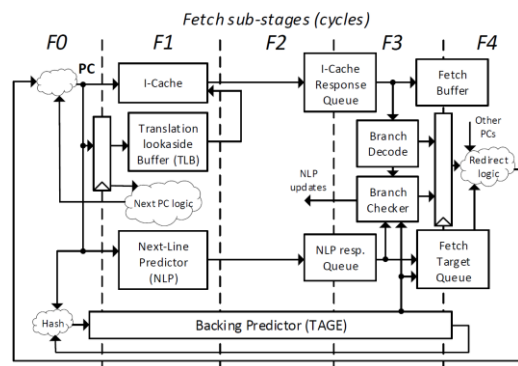
Outline

- ❑ SonicBOOM microarchitecture and its speculative execution vulnerabilities
- ❑ Spectre attacks and state of the art defenses
- ❑ SSE-RV architecture and mechanism
- ❑ Experimental results and analysis
- ❑ Future work

SonicBOOM microarchitecture

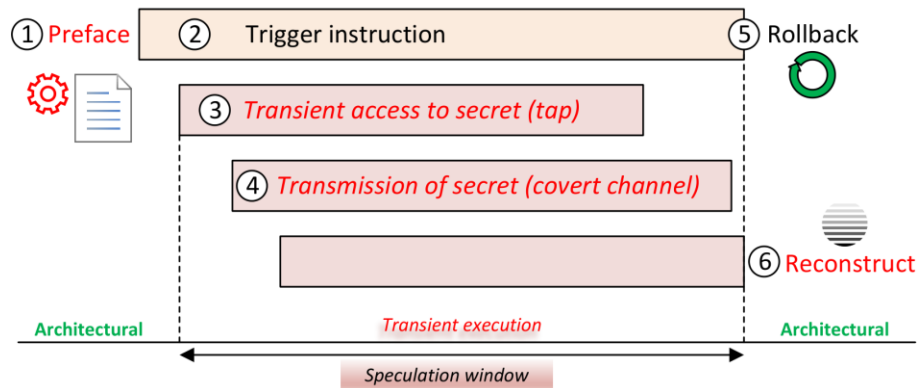
3rd-gen Berkeley Out-of-order Machine (SonicBOOM)[1]

- ❖ Out-of-order and speculative features allow for transient execution
- ❖ Comparable features to commercial processors



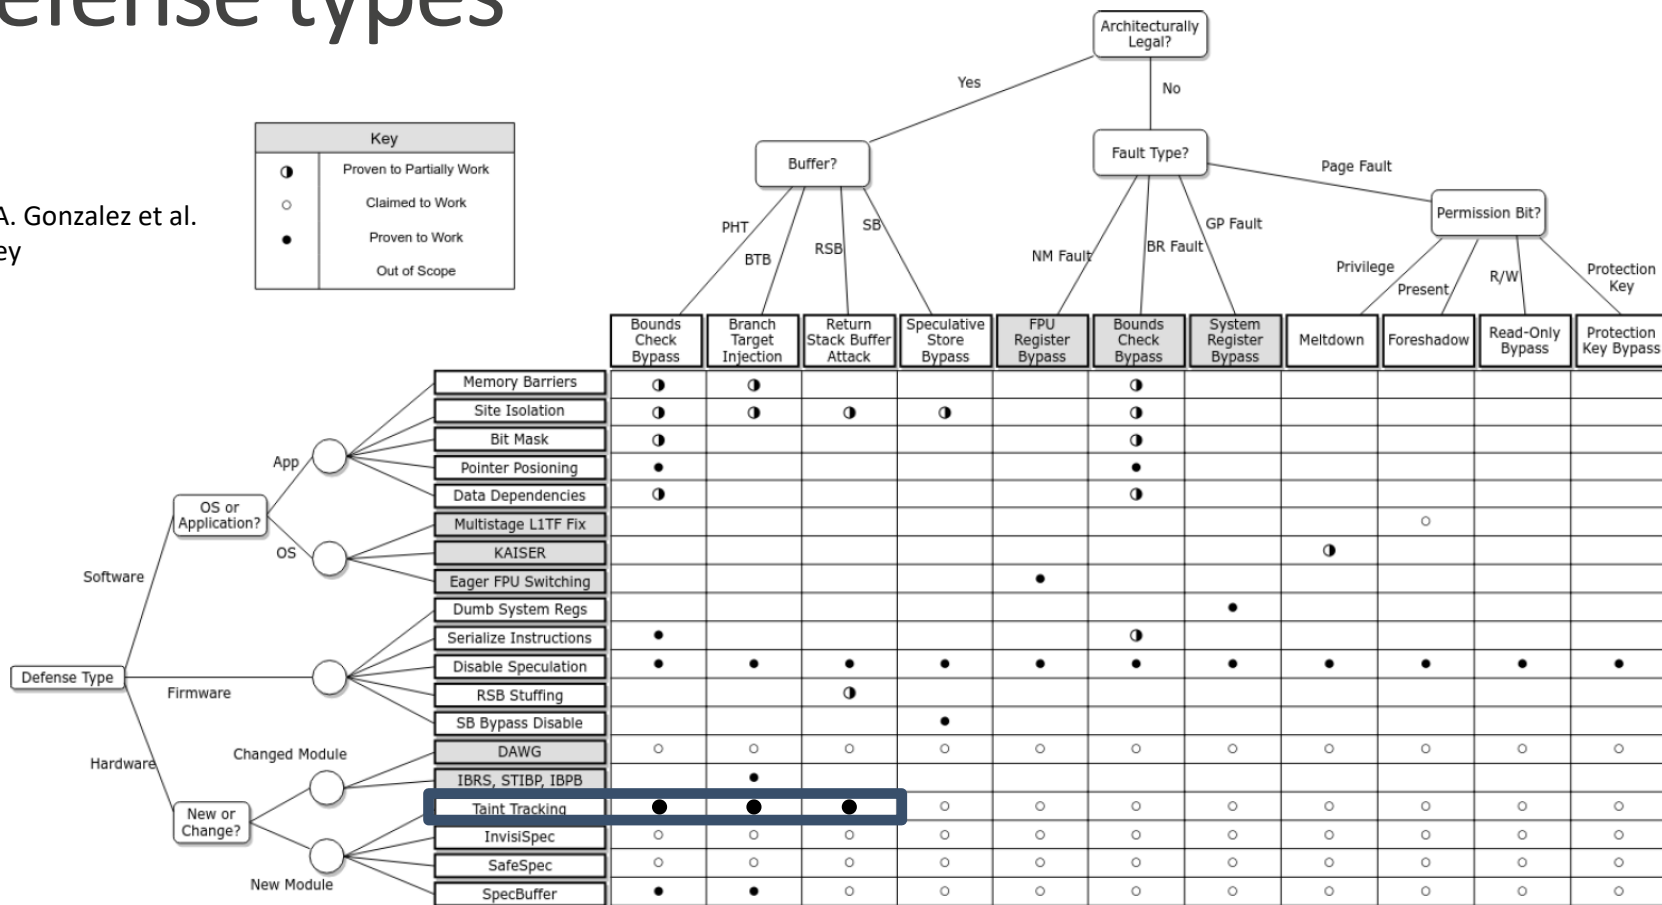
Spectre attacks

Transient execution attacks, including Spectre attacks, can be described in six major phases:



	Preface	Trigger		
		Instruction	Vulnerability	Attack agent
Spectre-v1	mis-train conditional branch direction prediction	conditional branch	out-of-bound offset	mis-speculated path
Spectre-v2	mis-train BTB for branch target prediction	call or jump	mis-speculated target	malicious function
Spectre-v5	craft a malicious gadget after a call site	return	RAS miss-match with software stack	gadget

Figure by A. Gonzalez et al.
UC Berkeley



State of the art defenses

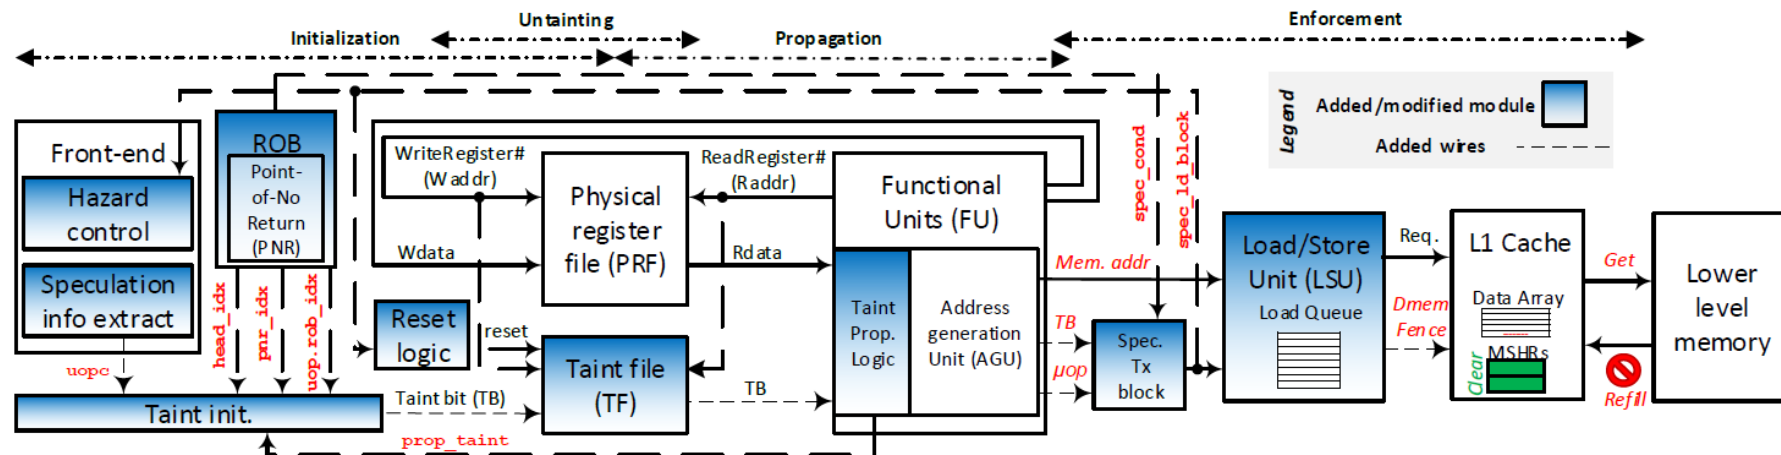
❑ Speculative Taint Tracking (STT)

- ❖ A **hardware-level taint tracking** scheme
- ❖ Taints data that is accessed during a speculative window and **delays** any subsequent **load** with tainted operands until they become untainted
- ❖ **Untainting** of registers happens **once the instruction that invoked it becomes non-speculative**

❑ Context-Sensitive Fencing (CSF)


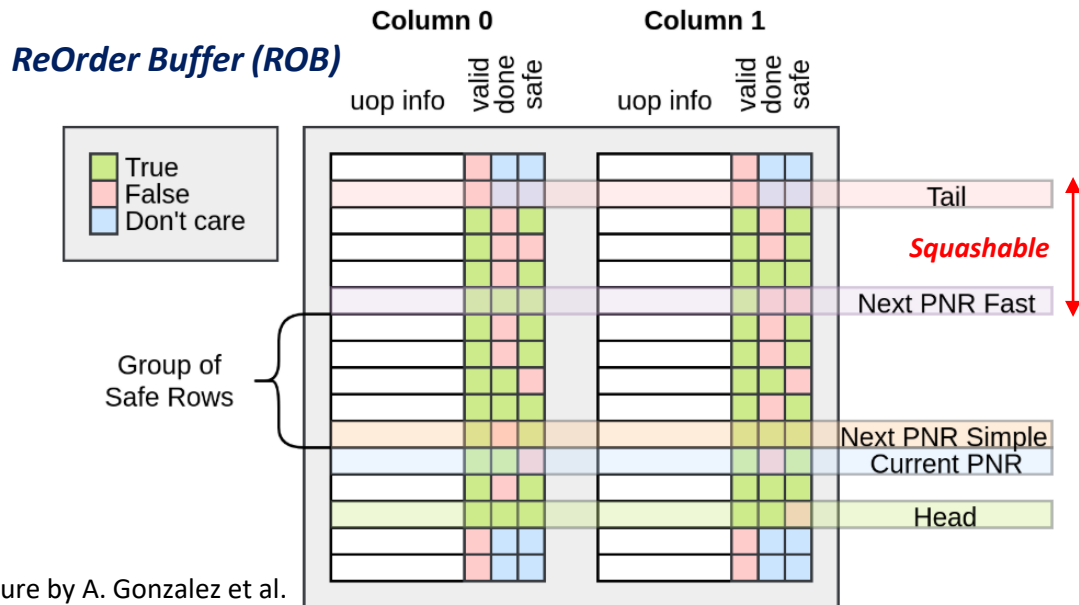
- ❖ A **micro-code level defense** against Spectre attacks
- ❖ **Dynamically alters the decoding** of the instruction stream, **injecting fences** only when dynamic conditions indicate they are needed

SSE-RV architectural overview



Point of No Return (PNR)

Speculation condition is determined by *fast* PNR[2]:




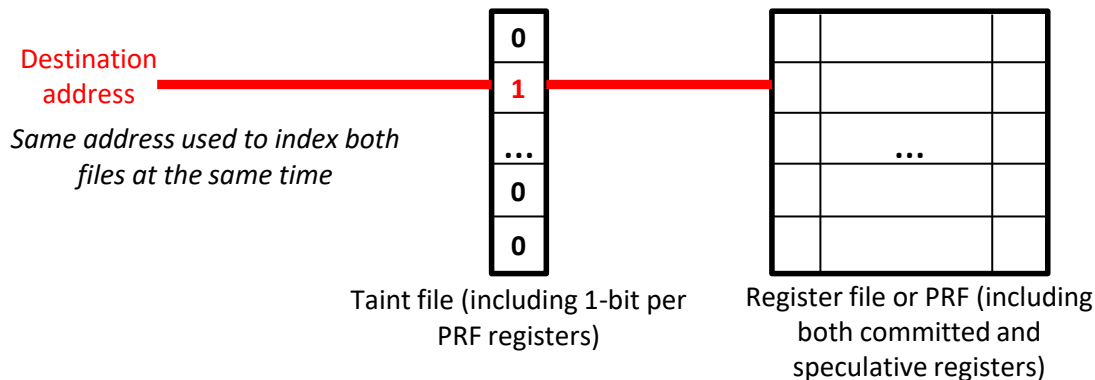
All forms of speculation is covered by PNR!

Figure by A. Gonzalez et al.
UC Berkeley

SSE-RV implementation - Initialization

Squashable (\approx speculative) **access** (load) instruction's **destination register**

- ❖ Squashable loads: the load is younger than PNR \rightarrow taint destination register
- ❖ Set the corresponding bit in taint file at the write-back stage for every squashable load



1

Initialization

2

Propagation

3

Enforcement

4

Untainting



SSE-RV implementation - Propagation

▪ Taints:

- ❖ Output of speculative access instructions
- ❖ Output of instructions with tainted inputs

▪ Propagation in parallel to functional units through 1-bit extended pipeline registers (chain of instructions) and synchronization queues

- ❖ Every in-flight instruction have #operands 1-bit taint inputs and a 1-bit taint output
 - $\text{Taint_out} = (\text{taint_in1} \mid \text{taint_in2} \mid \dots)$



1

Initialization

2

Propagation

3

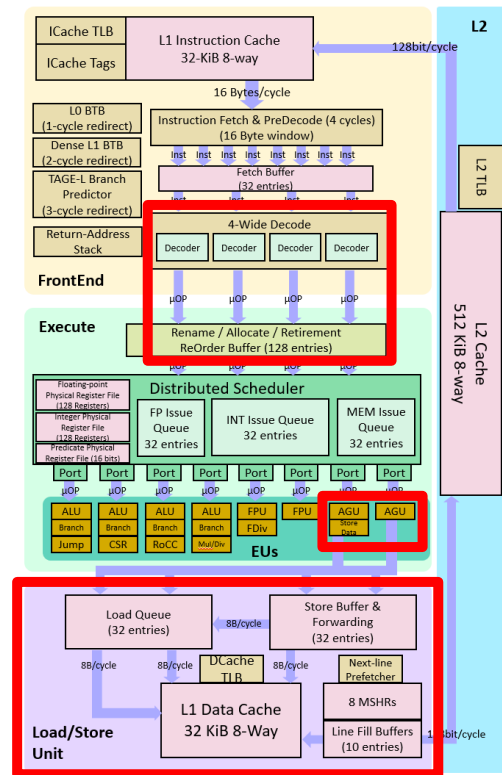
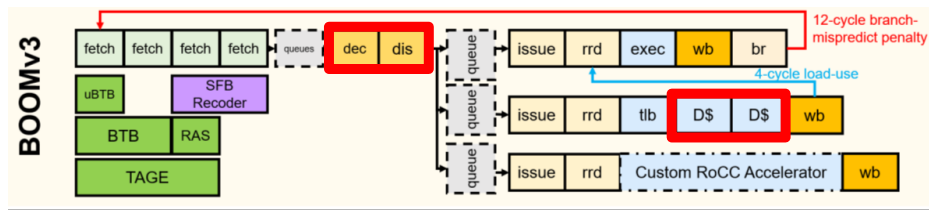
Enforcement

4

Untainting

SSE-RV implementation - Enforcement

- Detect transmit instructions (loads with tainted addresses) at the *memaddrcalc* (AGU) stage
- Delay load to cache until all operands are untainted (e.g., load address)
 - ❖ Block speculative cache refills
 - ❖ Backpressure on the dispatch units
 - Introduce a *dis* hazard to stall the pipeline, until the branch is resolved
 - ❖ Ops untainted → proceed with cache refills



1
Initialization

2
Propagation

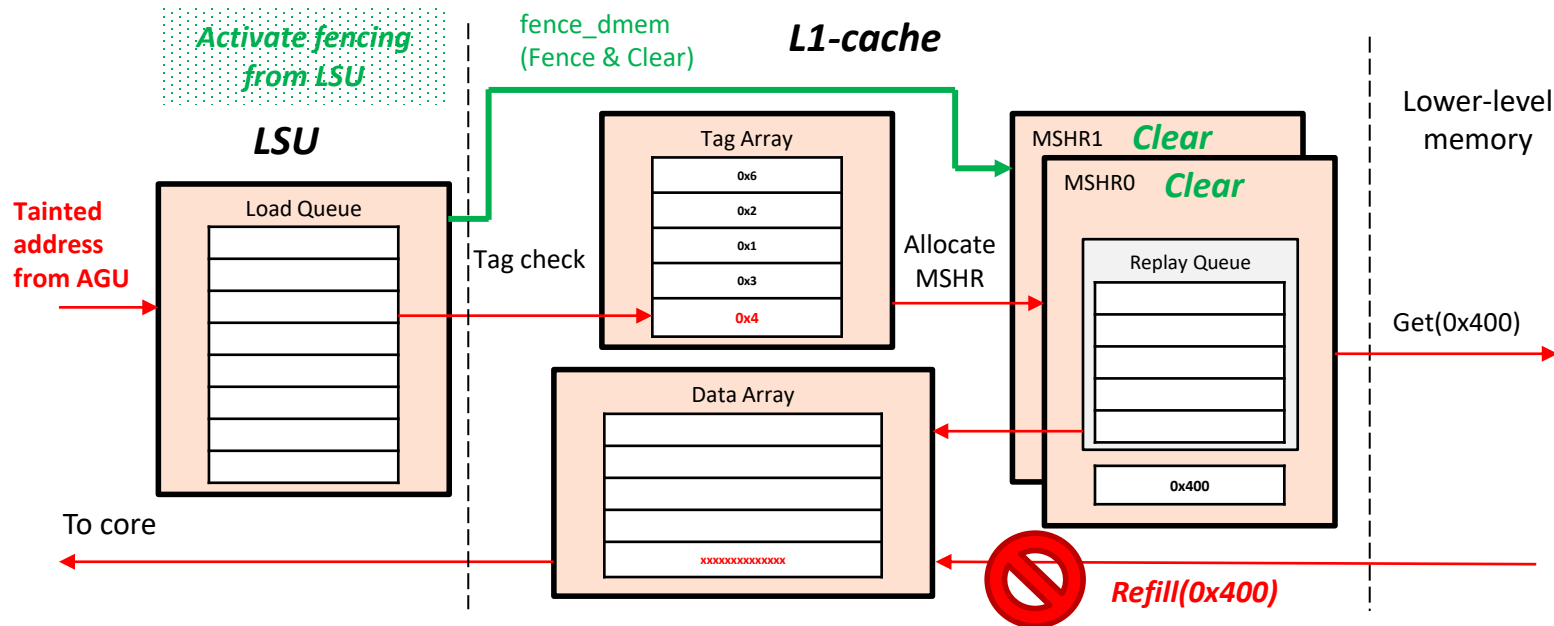
3
Enforcement

4
Untainting

SSE-RV implementation - Enforcement

■ Enable fencing to:

- ❖ Clear all the MSHRs and prefetch buffers
- ❖ Block the refill until the transmit instruction is no longer speculative



1

Initialization

2

Propagation

3

Enforcement

4


Untainting

SSE-RV implementation - Untainting

- **Disable the fencing and start untainting:**
 - ❖ A speculative access instruction becomes non-speculative (branch is resolved)
 - ❖ An instruction has all of its inputs untainted
- In case of misprediction, **the data and the corresponding taint become invalid** (squashed)
- In case of a correct prediction, the data is left intact, but **all the taint bits in the taint file will be reset to 0**
 - ❖ *Assuming there is only one batch of in-flight speculative instructions*



Taint file (including 1-bit per PRF registers)



Do not touch the data

1

Initialization

2

Propagation


3

Enforcement

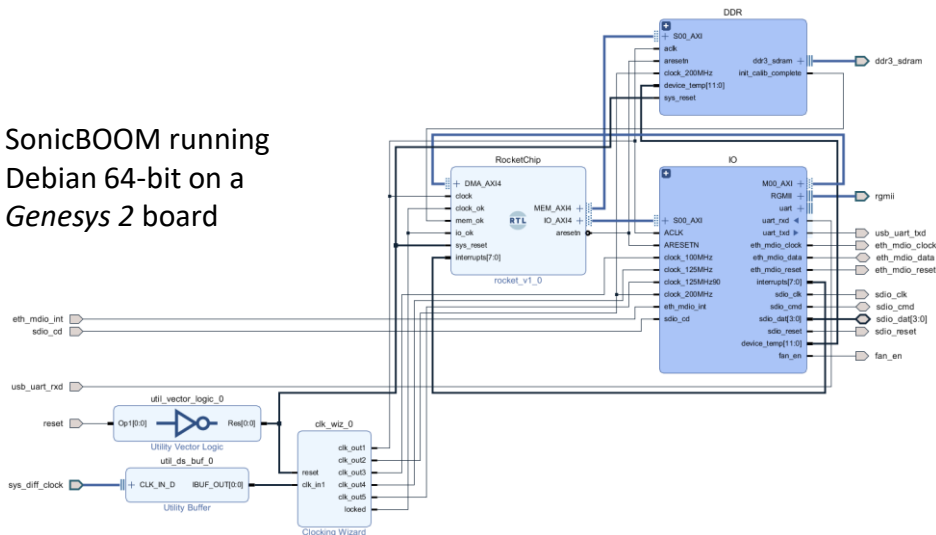
4

Untainting

Prototyping



SonicBOOM running
Debian 64-bit on a
Genesys 2 board



- ❑ Kintex-7 XC7K325T-2FFG900C on Genesys-2 board
- ❑ Core frequency = 100MHz
- ❑ Small core machine
 - ❖ 1-wide, 3-issue, 16KB D\$ (2 MSHRs), 16KB I\$, 512KB L2\$

SonicBOOM Spectre Attacks

- ❑ Replicated Spectre-v1, v2, and v5 on SonicBOOM and open sourced at:

https://chest.coe.neu.edu/?current_page=SOURCE_CODE

Attack	Cycles for one Byte	Bytes per Second (@100MHz)
Spectre-v1	4857203	20
Spectre-v2	4783403	21
Spectre-v5	196591	526

Protection results

Spectre-v1 on Original SonicBOOM

```

m[0x0x50328] = want(!) =?= guess(hits,dec,char) 1.(3, 33, !) 2.(1, 0, )
m[0x0x50329] = want(") =?= guess(hits,dec,char) 1.(2, 34, ") 2.(1, 0, )
m[0x0x5032a] = want(#) =?= guess(hits,dec,char) 1.(2, 35, #) 2.(1, 0, )
m[0x0x5032b] = want(S) =?= guess(hits,dec,char) 1.(3, 83, S) 2.(1, 0, )
m[0x0x5032c] = want(e) =?= guess(hits,dec,char) 1.(2, 101, e) 2.(1, 0, )
m[0x0x5032d] = want(c) =?= guess(hits,dec,char) 1.(4, 99, c) 2.(1, 0, )
m[0x0x5032e] = want(r) =?= guess(hits,dec,char) 1.(9, 114, r) 2.(1, 0, )
m[0x0x5032f] = want(e) =?= guess(hits,dec,char) 1.(2, 101, e) 2.(1, 0, )
m[0x0x50330] = want(t) =?= guess(hits,dec,char) 1.(2, 116, t) 2.(1, 0, )
m[0x0x50331] = want(I) =?= guess(hits,dec,char) 1.(6, 73, I) 2.(1, 0, )
m[0x0x50332] = want(n) =?= guess(hits,dec,char) 1.(2, 110, n) 2.(1, 0, )
m[0x0x50333] = want(T) =?= guess(hits,dec,char) 1.(7, 84, T) 2.(1, 0, )
m[0x0x50334] = want(h) =?= guess(hits,dec,char) 1.(6, 104, h) 2.(1, 0, )
m[0x0x50335] = want(e) =?= guess(hits,dec,char) 1.(2, 101, e) 2.(1, 0, )
m[0x0x50336] = want(S) =?= guess(hits,dec,char) 1.(4, 83, S) 2.(1, 0, )
m[0x0x50337] = want(o) =?= guess(hits,dec,char) 1.(1, 0, ) 2.(1, 111, o)
m[0x0x50338] = want(n) =?= guess(hits,dec,char) 1.(2, 110, n) 2.(1, 0, )
m[0x0x50339] = want(i) =?= guess(hits,dec,char) 1.(6, 105, i) 2.(1, 0, )
m[0x0x5033a] = want(c) =?= guess(hits,dec,char) 1.(2, 99, c) 2.(1, 0, )
m[0x0x5033b] = want(B) =?= guess(hits,dec,char) 1.(3, 66, B) 2.(1, 0, )
m[0x0x5033c] = want(O) =?= guess(hits,dec,char) 1.(1, 0, ) 2.(1, 79, 0)
m[0x0x5033d] = want(O) =?= guess(hits,dec,char) 1.(2, 79, 0) 2.(1, 0, )
m[0x0x5033e] = want(M) =?= guess(hits,dec,char) 1.(4, 77, M) 2.(1, 0, )

```

Secret extracted!

Spectre-v1 on SonicBOOM with SSE-RV

```

m[0x0x50388] = want(!) =?= guess(hits,dec,char) 1.(1, 0, ) 2.(0, 40, ( )
m[0x0x50389] = want(") =?= guess(hits,dec,char) 1.(1, 0, ) 2.(0, 0, )
m[0x0x5038a] = want(#) =?= guess(hits,dec,char) 1.(1, 0, ) 2.(0, 0, )
m[0x0x5038b] = want(S) =?= guess(hits,dec,char) 1.(1, 0, ) 2.(0, 0, )
m[0x0x5038c] = want(e) =?= guess(hits,dec,char) 1.(1, 0, ) 2.(1, 101, e)
m[0x0x5038d] = want(c) =?= guess(hits,dec,char) 1.(1, 0, ) 2.(0, 0, )
m[0x0x5038e] = want(r) =?= guess(hits,dec,char) 1.(1, 0, ) 2.(0, 0, )
m[0x0x5038f] = want(e) =?= guess(hits,dec,char) 1.(1, 0, ) 2.(0, 0, )
m[0x0x50390] = want(t) =?= guess(hits,dec,char) 1.(1, 0, ) 2.(0, 0, )
m[0x0x50391] = want(I) =?= guess(hits,dec,char) 1.(1, 0, ) 2.(0, 0, )
m[0x0x50392] = want(n) =?= guess(hits,dec,char) 1.(1, 0, ) 2.(0, 0, )
m[0x0x50393] = want(T) =?= guess(hits,dec,char) 1.(1, 0, ) 2.(0, 0, )
m[0x0x50394] = want(h) =?= guess(hits,dec,char) 1.(1, 0, ) 2.(0, 0, )
m[0x0x50395] = want(e) =?= guess(hits,dec,char) 1.(1, 0, ) 2.(0, 0, )
m[0x0x50396] = want(S) =?= guess(hits,dec,char) 1.(1, 0, ) 2.(0, 0, )
m[0x0x50397] = want(o) =?= guess(hits,dec,char) 1.(1, 0, ) 2.(0, 0, )
m[0x0x50398] = want(n) =?= guess(hits,dec,char) 1.(1, 0, ) 2.(0, 0, )
m[0x0x50399] = want(i) =?= guess(hits,dec,char) 1.(1, 0, ) 2.(0, 0, )
m[0x0x5039a] = want(c) =?= guess(hits,dec,char) 1.(1, 0, ) 2.(0, 0, )
m[0x0x5039b] = want(B) =?= guess(hits,dec,char) 1.(1, 0, ) 2.(0, 0, )
m[0x0x5039c] = want(O) =?= guess(hits,dec,char) 1.(1, 0, ) 2.(0, 0, )
m[0x0x5039d] = want(O) =?= guess(hits,dec,char) 1.(1, 0, ) 2.(1, 79, 0)
m[0x0x5039e] = want(M) =?= guess(hits,dec,char) 1.(1, 0, ) 2.(0, 0, )


```

Secret hidden!

Similar results for Spectre-v2 and v5

Performance impact

Configuration	CoreMark Size	Total cycles	Total time (secs)	Iterations/Sec (CoreMark)	CoreMark/MHz	Iterations
Original small core SonicBoom machine	666	43,127	43.1	231.9	2.32	10,000
SSE-RV protected small core SonicBoom machine	666	153,992	154.0	64.9	0.65	10,000



The MiBench[3] performance overhead
when only protecting against Spectre-v1

More transmitter instructions → more overhead!

[3] M.R. Guthaus, J.S. Ringenberg, D. Ernst, T.M. Austin, T. Mudge, and R.B. Brown, MiBench: A free, commercially representative embedded benchmark suite, WWC-4

Implementation Cost

Unprotected vs. SSE-RV FPGA resources utilization

Configuration	LUT	FF	BRAM	DSP
Unprotected	109790	72930	175	36
SSE-RV	112549(+2.5%)	73414(+0.7%)	175	36

Unprotected vs. SSE-RV core area estimates, 130nm BiCMOS technology

Configuration	Gate count	Total area (μm^2)	Area scaling
Unprotected	509,547	3913320.96	1
SSE-RV	511,693	3929796.48	1.0042

Unprotected vs. SSE-RV core power estimates, 130nm BiCMOS technology

Configuration	Leakage (mW)	Dynamic (mW)	Total (mW)
unprotected	0.2065	286.12	286.33
SSE-RV	0.2071	291.24	291.44

Future work

- ❑ **Enhance the SSE-RV performance** in both single and multi-core settings
- ❑ **Explore the vulnerability** of SonicBOOM to **MDS attacks**
- ❑ **Extend** our methodology to defend **against software attacks**

Thank you!

Comments/Questions?

sabbagh.m@northeastern.edu

This work was supported in part by the National Science Foundation under Grants SaTC1563697, CNS-1916762, and with the industry support from the Center for Hardware and Embedded Systems Security and Trust (CHEST) and Draper Laboratory.