

# rtl\_v: push-button verification of software on hardware

**CARRV 2021**

Noah Moroze  
Anish Athalye  
Frans Kaashoek  
Nickolai Zeldovich

MIT CSAIL

# Formal verification

- *Proves* that a system satisfies a certain property
  - Distinct from traditional “verification” or testing
- Increases confidence in security and correctness

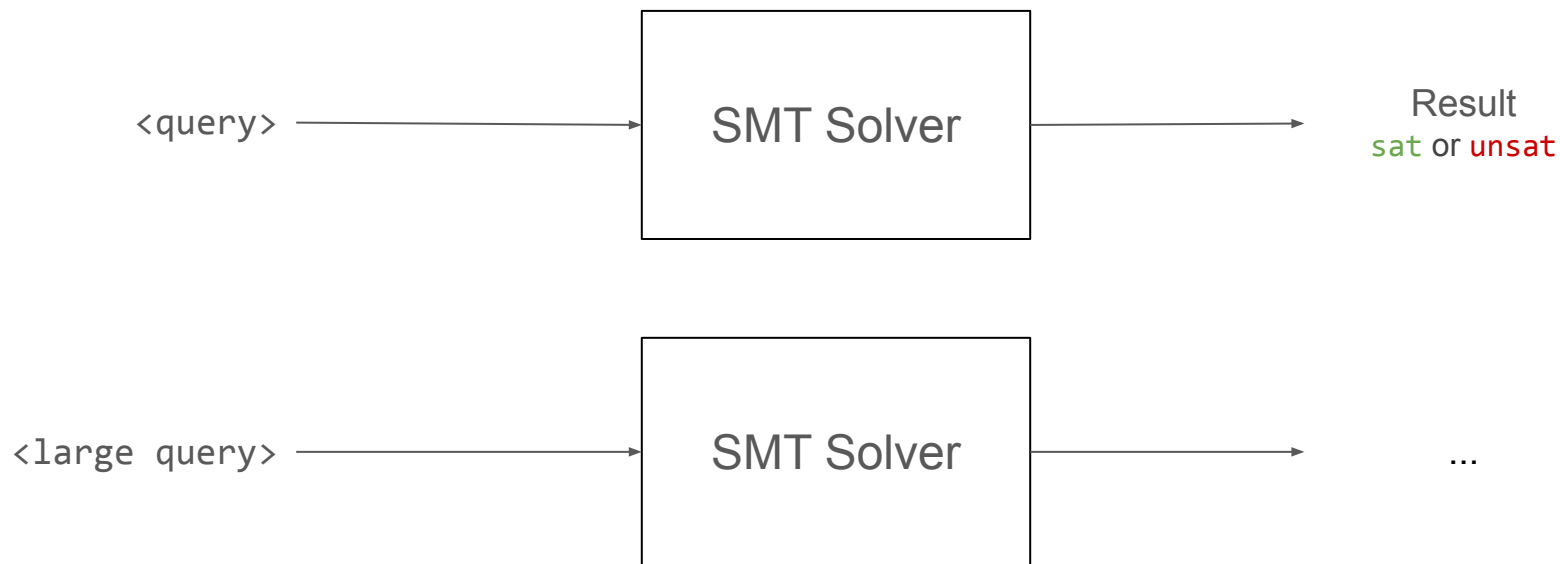
# Current verification approach

- SymbiYosys is a popular open-source hardware formal verification tool
- Developers express properties using SystemVerilog Assertion (SVA) constructs

```
module counter (  
    input clk,  
    input rst_n,  
    output reg [1:0] out  
);  
  
    always @(posedge clk) begin  
        if (!rst_n || out == 2'd2) begin  
            out <= 2'd0;  
        end else begin  
            out <= out + 1'b1;  
        end  
    end  
  
    assert(out != 2'd3);  
  
endmodule
```

# SMT-based formal verification

- Circuit logic and property encoded into SMT query, evaluated by SMT solver
- Complex queries can be a performance bottleneck
  - Solver may take a long time, or time out completely!



# Limitation: reasoning about software on hardware

- Want to prove that a CPU's state is cleared by boot code (*deterministic start*)
- Requires reasoning about execution of software on hardware
- SymbiYosys inherently produces complex SMT queries when reasoning about long executions
- Cannot be done by existing tools: we need a new approach!

# rtl\_v

- rtl\_v is an approach for verifying circuits that enables reasoning about many cycles of execution
- It achieves this by:
  - Using symbolic execution to generate more efficient verification queries
  - Supporting *circuit-agnostic property checkers* that enable disciplined performance optimizations

# Generating efficient verification queries

# SymbiYosys' SMT query generation

```
module foo (  
    input clk,  
    input rst_n,  
    output reg out  
);  
  
always @(posedge clk) begin  
    if (!rst_n) begin  
        out <= 1'b0;  
    end else begin  
        out <= ~out;  
    end  
end  
  
endmodule
```

Yosys

```
foo_t(state, next_state):  
    return  
    next_state.out ==  
    (rst_n ? !state.out : 0)
```

**Verilog**

**SMT-LIB Relation**



# SymbiYosys' SMT query generation


- Want to prove that a property  $P(s)$  holds for some state  $s$
- Construct  $n$  symbolic circuit states  $s_0, s_1, \dots, s_n$
- Constrain  $s_0$  to be a possible initial state ( $I(s_0)$  is true)
- Can use transition relation  $T(s_a, s_b)$  to encode execution in query

$\text{assert}(I(s_0) \wedge T(s_0, s_1) \wedge T(s_1, s_2) \wedge \dots \wedge T(s_{n-1}, s_n) \wedge \neg P(s_n))$

---

SMT solvers are bad at reasoning  
about unrolled complex execution

# rtlv approach



**SMT-LIB relation**

**Step function**

# Generating SMT query through symbolic execution

`assume(I( $s_0$ ))`

`$s_1$  = step( $s_0$ , rst=1)`

`$s_2$  = step( $s_1$ , rst=0)`

`...`

`$s_n$  = step( $s_{n-1}$ , rst=0)`

`assert( $\neg P(s_n)$ )`

---

Execution modelled by  
Rosette's symbolic  
execution system

Query does not encode execution logic, and  
therefore does not inherently scale with  
execution length


# Good applications

- Although final query no longer inherently scales linearly with  $n$ , query is only fast if  $s_n$  is small
- Some applications are better for this:
  - Lots of concrete state (e.g. known code with concrete control flow)
  - Deterministic start is a great example!

# SymbiYosys SMT solver call times out

- Experiment: verify deterministic start for PicoRV32 CPU with SymbiYosys and rtlv
- Boot code takes 104 cycles to execute
- SymbiYosys times out after 12 hours, rtlv finishes in 1.3 seconds
- Problem goes from intractable to fast!

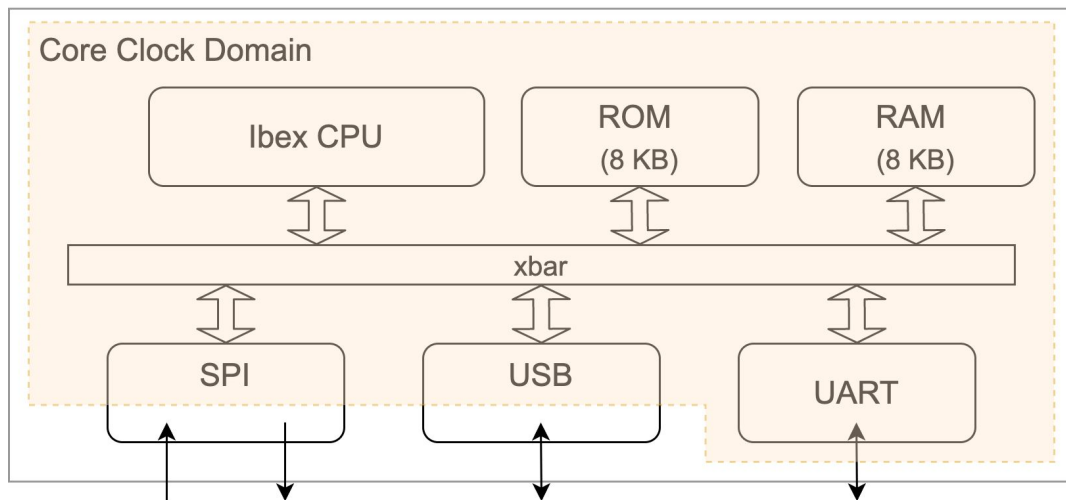
Partial deterministic  
start experiment



# Optimizing verification of complex hardware using property checkers

# Case Study: MicroTitan

- Subset of OpenTitan RISC-V SoC
- Hardware complexity:  
4,300 FFs (~3x PicoRV32)
- Boot code length:  
24,516 cycles (~235x)



# Scaling problem: SPI state machine

Cycles 0 - 149

Cycle 150

State machine converges

Symbolic external inputs

spi.  
rxfr\_ctrl\_st

Symbolic state  
feeds back and  
causes term to  
grow

No optimization

spi.  
rxfr\_ctrl\_st

Symbolic term  
continues to  
grow  $\Rightarrow$   
intractable final  
query

Make  
spi.rxfr\_ctrl\_st  
concrete

spi.  
rxfr\_ctrl\_st  
IDLE

Register  
remains  
concrete  $\Rightarrow$   
fast final  
solver query



# Circuit-agnostic property checkers



- rtlv/shiva is *sound*: changing the list of hints cannot make the verification tool incorrectly return 'true'
- Ensured through auxiliary solver queries
  - E.g. concretize hint uses SMT solver to determine if field evaluates to concrete value
- Existing tools do not support this approach, since they don't provide direct access to solver

# Results

- Verification runtime: 100 minutes
- Violations found:
  - SPI TX leak
  - SPI RX leak
  - USB buffer leak
  - Reset line leak
- Contributed back to OpenTitan:  
<https://github.com/lowRISC/opentitan/pull/2420>

# Limitations

- Circuit features
  - No latches
  - No async structures
- Complexity of HW/SW
  - Expect we can verify 100s of lines of C on microcontroller
  - Do not expect that current system can verify full OS

# Summary

- `rtl_v` is an approach for verifying long executions on hardware
- Two key ideas:
  - Symbolically execute circuit using Rosette to generate efficient solver queries
  - Develop circuit-agnostic property checkers to optimize and scale verification in a disciplined way
- `rtl_v` enables you to verify properties that can't otherwise be verified using open and accessible tools!

# Thank you!

- PicoRV32 comparison case study:  
<https://github.com/anishathalye/deterministic-start-benchmark>
- MicroTitan case study: <https://github.com/nmoroze/kronos>
  - Blog post: <https://noahmoroze.com/kronos.html>
  - Thesis: <https://pdos.csail.mit.edu/papers/moroze-meng.pdf>